

Kai Hwang

*Cloud Computing for Machine
Learning and Cognitive Applications*

The MIT Press, Cambridge, MA, 2017

Chapter 3: Virtual Machines, Containers and Cluster Operating Systems

(Slides for 4 -hour lectures)

All rights reserved by Kai Hwang and MIT Press, 2017.

For exclusive use by qualified instructors adopting
the textbook. Not for commercial or publication release.

Difference between traditional computer and virtual machines

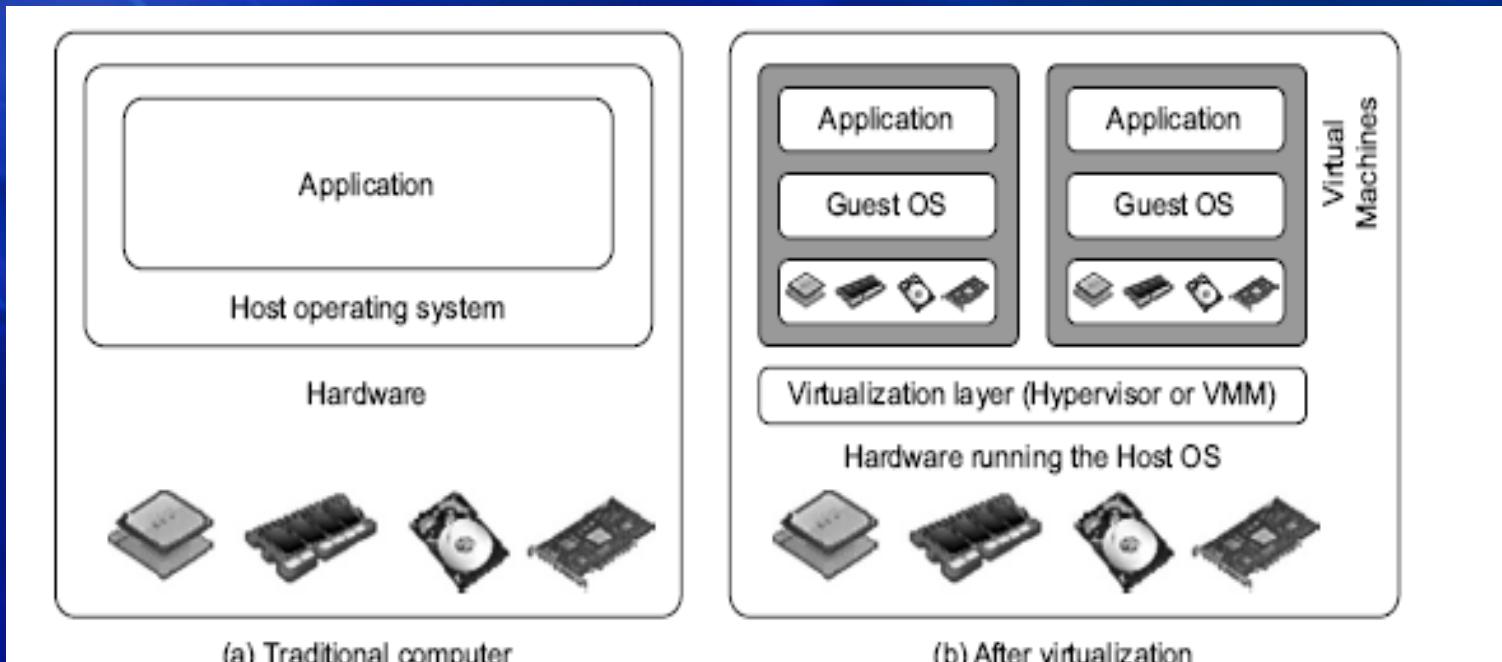


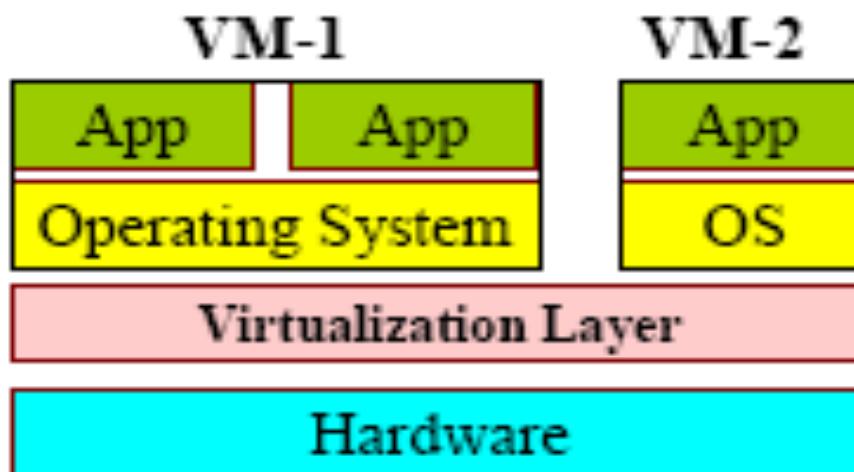
FIGURE 3.1

The architecture of a computer system before and after virtualization, where VMM stands for virtual machine monitor.

(Courtesy of VMWare, 2008)

What is virtualization?

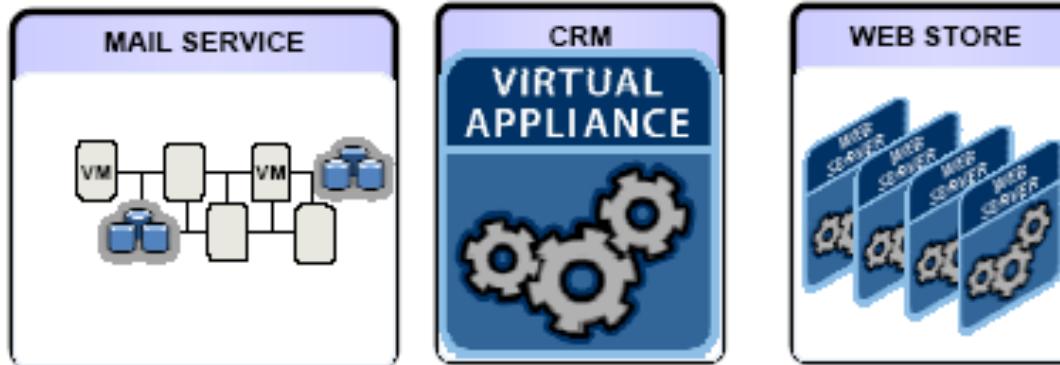
- A level of indirection between hardware and software.



- Virtual Machine abstraction
 - Run all software written for physical machine.

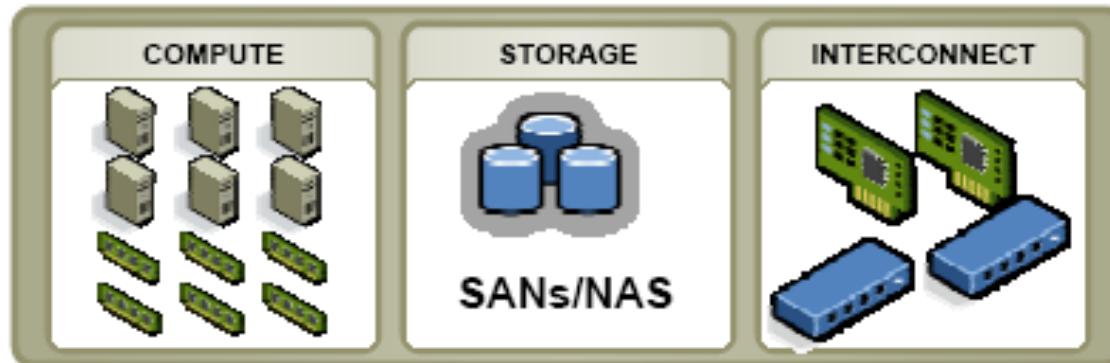
User's view of virtualization

LOGICAL VIEW



Virtualization Layer - Optimize HW utilization, power, etc.

PHYSICAL VIEW



(Courtesy of VMWare, 2008)

Various interface abstraction levels in modern computer architecture

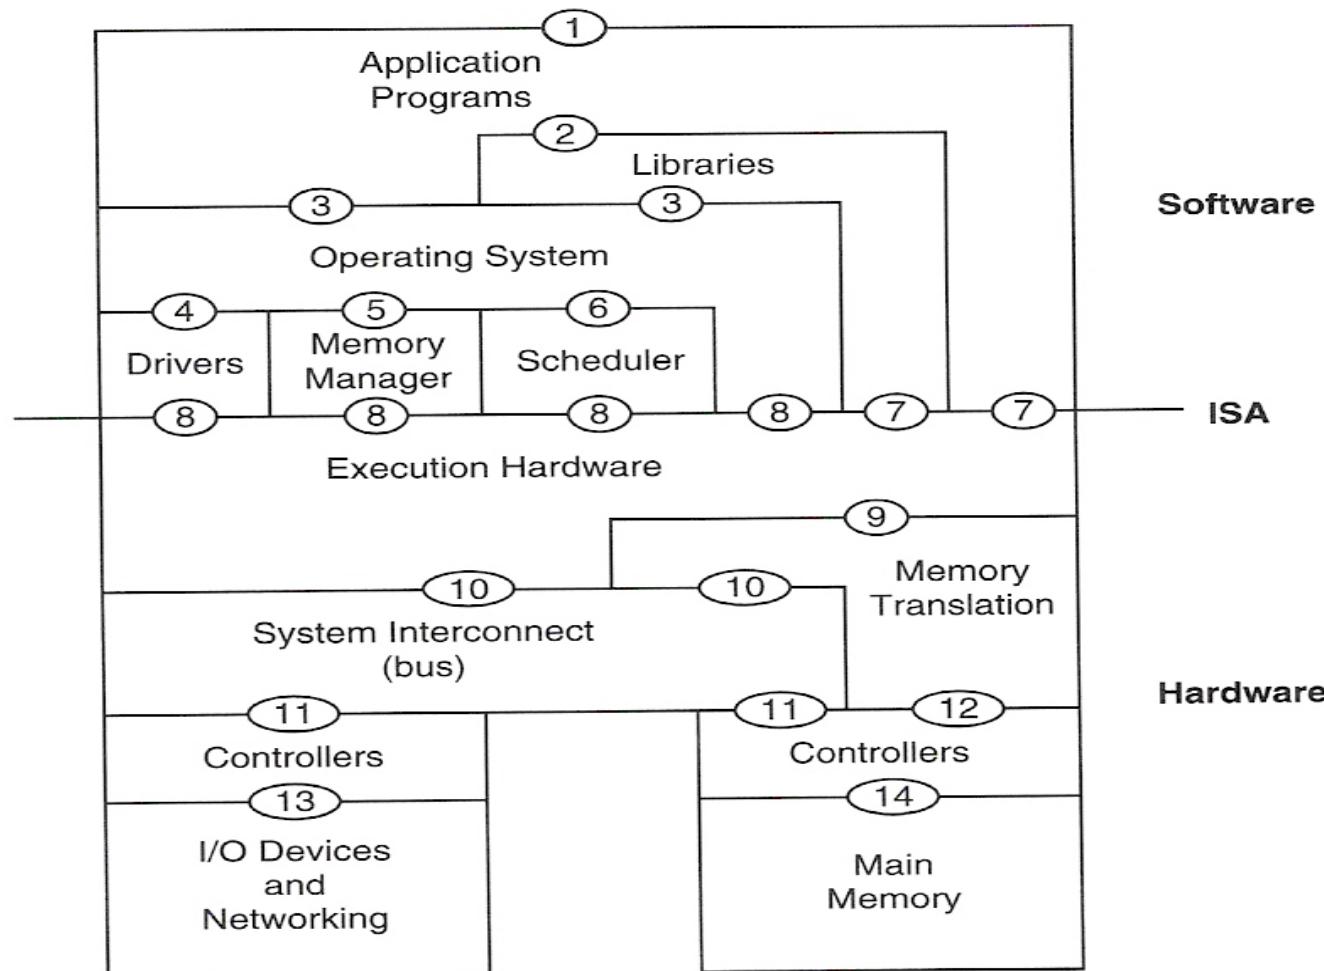
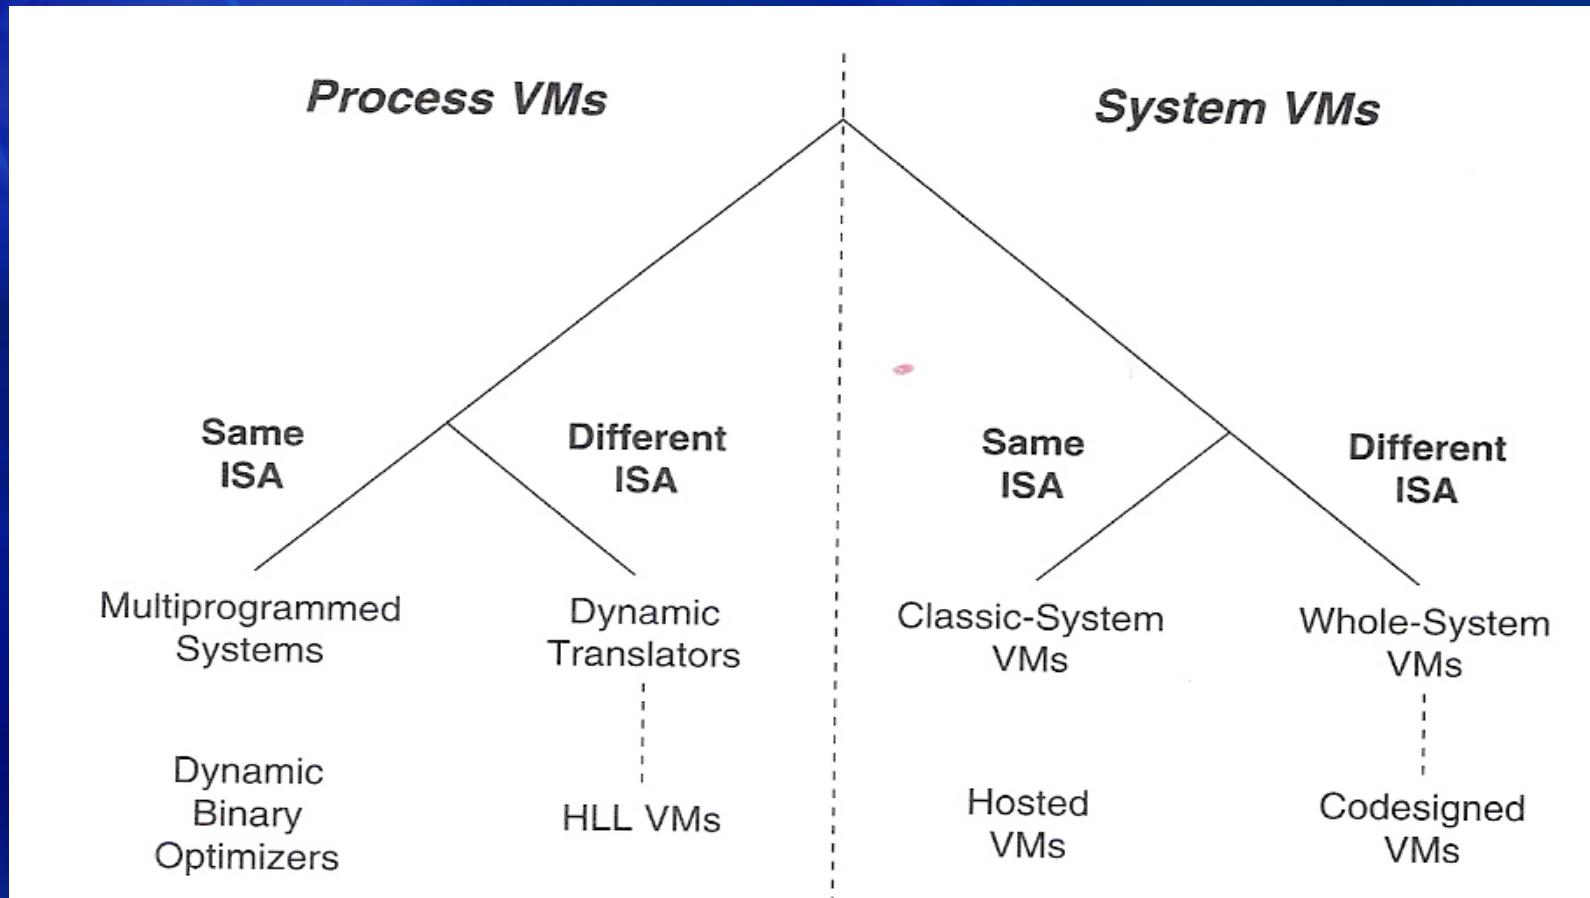
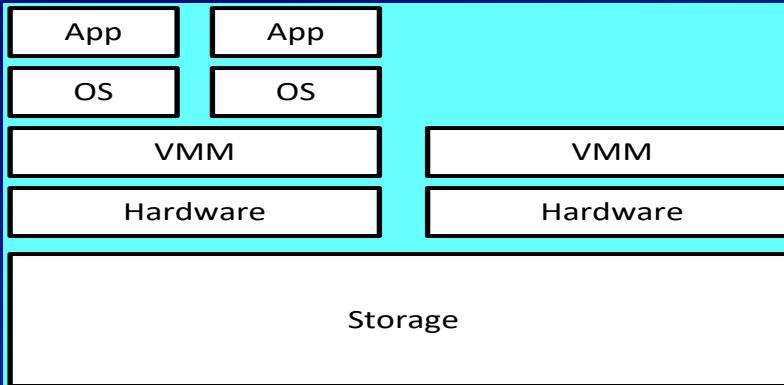


Figure 1.4 Computer System Architectures. *Implementation layers communicate vertically via the shown interfaces. This view of architecture is styled after one given by Glenford Myers (1982).*

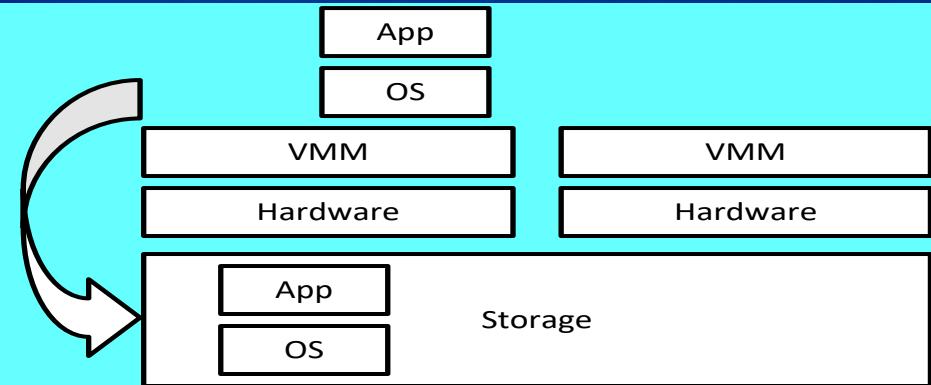
A taxonomy of virtual machines



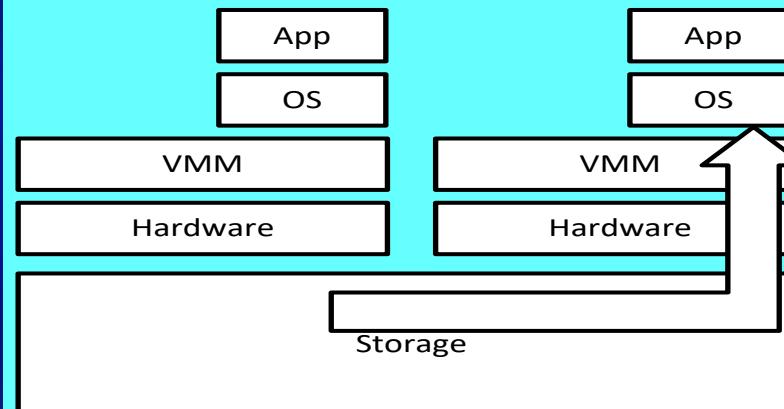
(Courtesy of Smith and Nail, 2005)



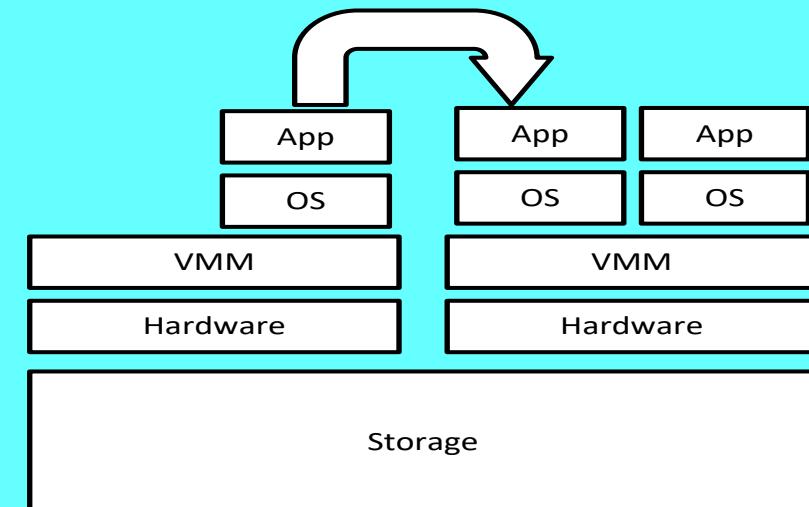
(a) Multiplexing



(b) Suspension(Storage)



(c) Provision(Resume)



(d) Life Migration

Figure 3.2 Virtual machine multiplexing, suspension, provision, and migration in a distributed computing environment

Virtualization at various abstraction levels

Table 1.7 Relative Merits of Virtualization at Five Abstraction Levels

Level of Virtualization	Functional Description	Example Pakages	Relative Merits, App Flexibility/Isolation, and Implementation Complexity
Instruction Set Architecture	<i>Emulation of a guest ISA by the host ISA</i>	Dynamo, Bird, Bochs, Crusoe	<i>Very Low performance, high app flexibility, and median complexity and Isolation</i>
Hardware-level Virtualization	<i>Virtualization on top of bare metal hardware</i>	XEN, VMWare, Virtusl PC	<i>High performance and complexity, median app flexibility, and good app isolation</i>
Operating System Level	<i>Isolated containers as OS instances</i>	Docker Engine, Jail, FVM,	<i>Highest performance, Low App flexibility and Isolation, and average complexity</i>
Run-Time Library Level	<i>Creating VM via run-time library through API hooks</i>	Wine, cCUDA, WABI, LxRun	<i>Average performance, low app flexibility and isolation, and low complexity</i>
User Application Level	<i>Deploy HLL VMs at user app level</i>	JVM, .NET CLR, Panot	<i>Low performance and app flexibility, very high complexity and app isolation</i>

Relative performance of virtualization at 5 abstraction levels

Table 3.1 Relative Merits of Virtualization at Various Levels

Level of Implementation	Higher Performance	Application Flexibility	Implementation Complexity	Application Isolation
ISA	X	XXXXX	XXX	XXX
Hardware-level virtualization	XXXXX	XXX	XXXXX	XXXX
OS-level virtualization	XXXXX	XX	XXX	XX
Runtime library support	XXX	XX	XX	XX
User application level	XX	XX	XXXXX	XXXXX

Virtualization at ISA level

Emulating a given ISA by the ISA of the host machine. For example, MIPS binary code can run on an x-86-based host machine with the help of an ISA emulator. Example systems: Bochs, Crusoe, Qemu, BIRD, Dynamo

Advantage: It can run a large amount of legacy binary codes written for various processors on any given new hardware host machines. Best application flexibility

Shortcoming & limitation: One source instruction may require tens or hundreds of native target instructions to emulate, which is rather slow. The V-ISA requires adding a processor-specific software translation layer to do the emulation aided by the compiler

Virtualization at hardware abstraction level

This is known as **hypervisor**. Virtualization is performed right on top of the bare metal hardware. It generates virtual hardware environments for VMs.

Example systems: Xen, Hyper V, KVM, Virtual PC, Denali

Advantage: it has higher performance and good application isolation

Shortcoming and limitations: very expensive to implement due to its complexity

Full virtualization vs. para-virtualization

Full virtualization does not need to modify guest OS, and critical instructions are emulated by software through the use of binary translation. On the other hand, para-virtualization needs to modify guest OS, and non-virtualizable instructions are replaced by hypercalls that communicate directly with the hypervisor or VMM. As of full virtualization, the advantage is not having to modify OS. However, this approach of binary translation slows down the performance considerably

Para-virtualization reduces the overhead, but the cost of maintaining para-virtualized OS is high. The improvement depends on the workload. VMware Workstation applies full virtualization, which uses binary translation to automatically modify x86 software on-the-fly to replace critical instructions. The para-virtualization is supported by Xen, Denali and VMware ESX

Full virtualization

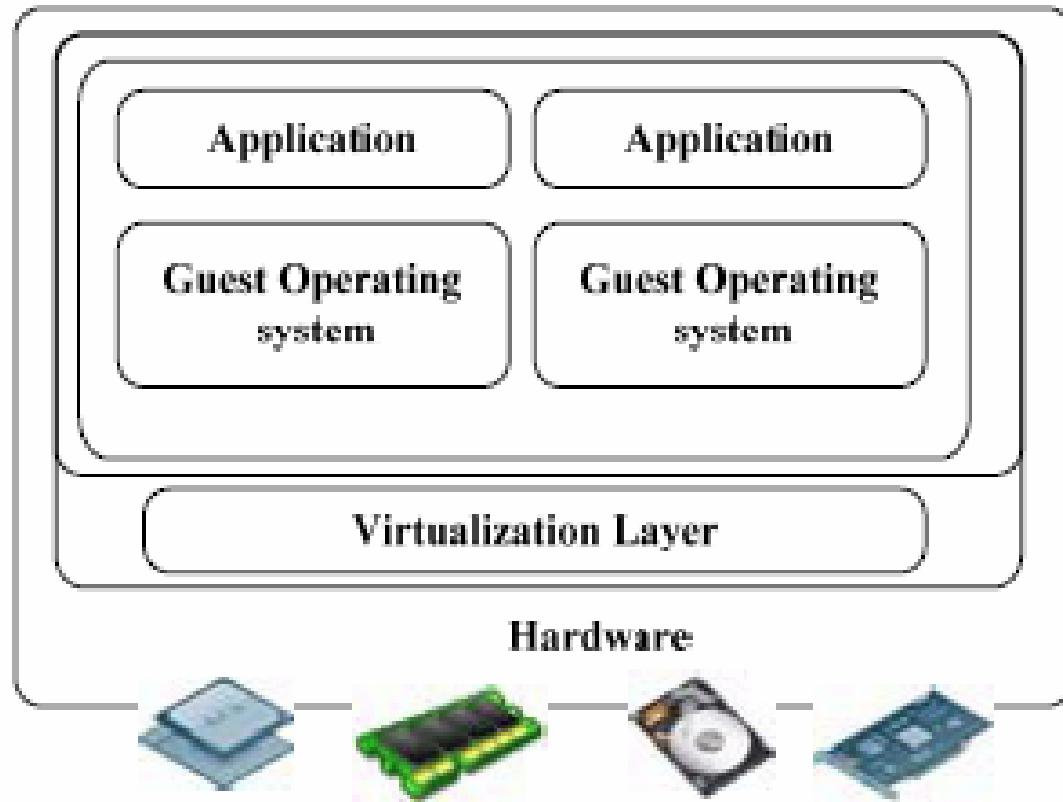


Figure 6.9 The concept of full virtualization using a hypervisor or a VMM directly sitting on top of the bare hardware devices. Note that no host OS is used here as in Figure 6.11.

Hypervisor

A hypervisor is a hardware virtualization technique allowing multiple operating systems, called guests, to run on a host machine. This is also called the Virtual Machine Monitor (VMM).

Type 1 hypervisor or the bare metal hypervisor sits on the bare metal computer hardware like the CPU, memory, etc. All the guest operating systems are a layer above the hypervisor. The hypervisor is the first layer over the hardware, such as the original CP/CMS hypervisor developed by IBM. An example is Microsoft Hyper-V

Type 2 or the hosted hypervisor does not run over the bare metal hardware, but over a host operating system. The hypervisor is the second layer over the hardware. The guest operating systems runs a layer over the hypervisor and form the third layer. An example is FreeBSD. The operating system is usually unaware of the virtualization

Virtual machine architectures

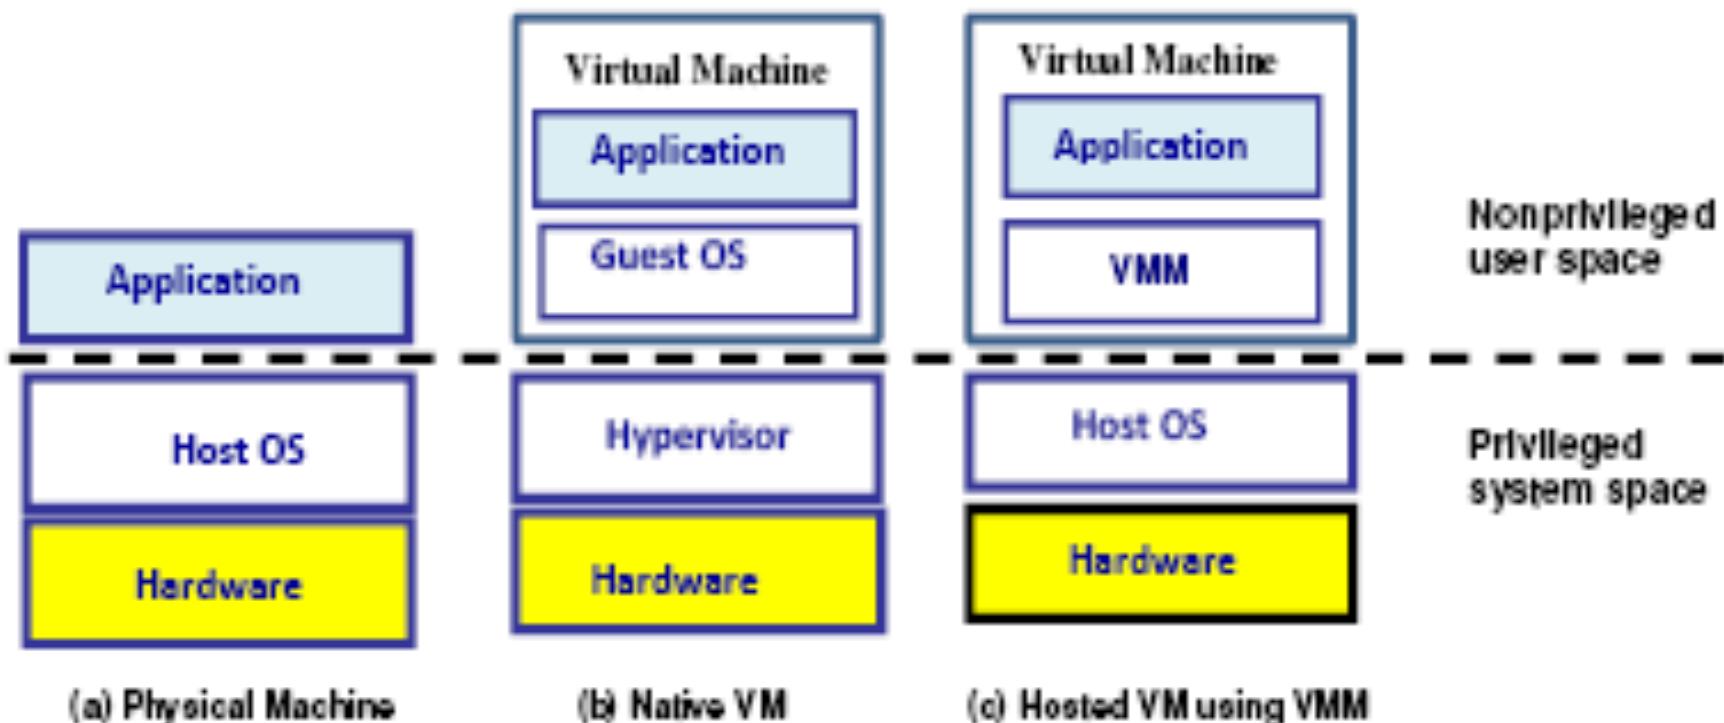


Figure 1.20 Three major components of a conventional physical machine is shown in Part (a). Native VM shown in Part (b) is created by a hypervisor sitting on top of bare metal hardware. A hosted VM shown in Part (c) is created by a VMM middleware in cooperation with the host OS.

Table 3.4

Hypervisors or VM monitors for generating VMs

Hypervisor	Host CPU	Host OS	Guest OS	Architecture, Applications, and User Community
XEN	x86, x86-64, IA-64	NetBSD, Linux	Linux, Windows, BSD, Solaris	Native hypervisor (Example 3.1) developed at Cambridge University
KVM	x86, x86-64, IA-64, S-390, PowerPC	Linux	Linux, Windows, FreeBSD, Solaris	Hosted hypervisor based on paravirtualization at the user space
Hyper V	x86 based	Server 2003	Windows servers	Windows-based native hypervisor, marketed by Microsoft
VMWare Workstation, VirtualBox	x86, x86-64	Any host OS	Windows, Linux, Darwin Solaris, OS/2, FreeBSD	Hosted hypervisor with a paravirtualization architecture

Hypervisor and the XEN architecture

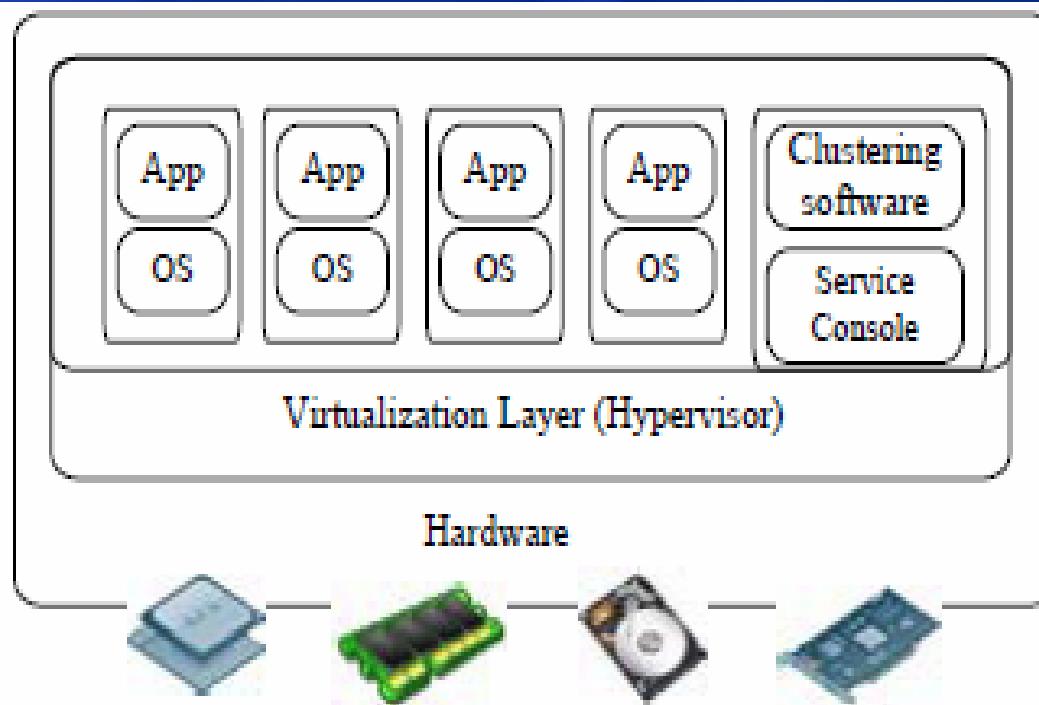


Figure 6.7 A hypervisor is the software layer for virtualization of the bare metal hardware. This layer can be implemented as a micro-kernel of the OS. It converts physical devices into virtual resources for user applications to run on the virtual machines deployed.

The XEN architecture (1)

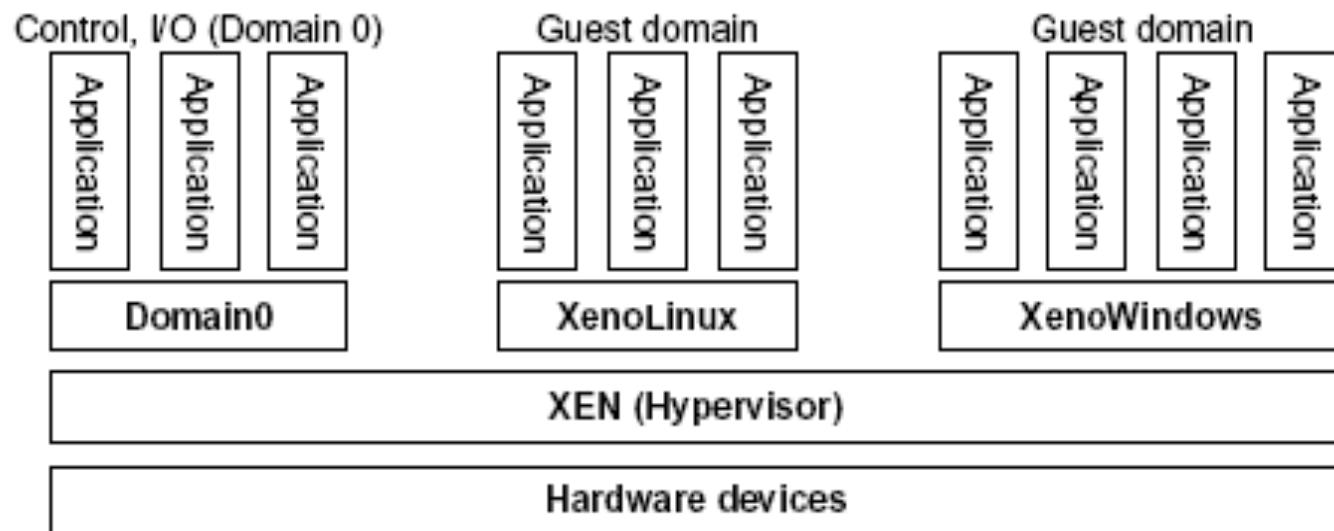


FIGURE 3.5

The Xen architecture's special domain 0 for control and I/O, and several guest domains for user applications.

The XEN architecture (2)

Xen is an open source hypervisor program developed by Cambridge University. Xen is a micro-kernel hypervisor, which separates the policy from the mechanism. The Xen hypervisor implements all the mechanisms, leaving the policy to be handled by Domain 0, as shown in Figure 3.5. Xen does not include any device drivers natively [7]. It just provides a mechanism by which a guest OS can have direct access to the physical devices. As a result, the size of the Xen hypervisor is kept rather small. Xen provides a virtual environment located between the hardware and the OS. A number of vendors are in the process of developing commercial Xen hypervisors, among them are Citrix XenServer [62] and Oracle VM [42].

The XEN Architecture (2)

Xen is an open source hypervisor program developed by Cambridge University. Xen is a micro-kernel hypervisor, which separates the policy from the mechanism. The Xen hypervisor implements all the mechanisms, leaving the policy to be handled by Domain 0, as shown in Figure 3.5. Xen does not include any device drivers natively [7]. It just provides a mechanism by which a guest OS can have direct access to the physical devices. As a result, the size of the Xen hypervisor is kept rather small. Xen provides a virtual environment located between the hardware and the OS. A number of vendors are in the process of developing commercial Xen hypervisors, among them are Citrix XenServer [62] and Oracle VM [42].

Host-based virtualization

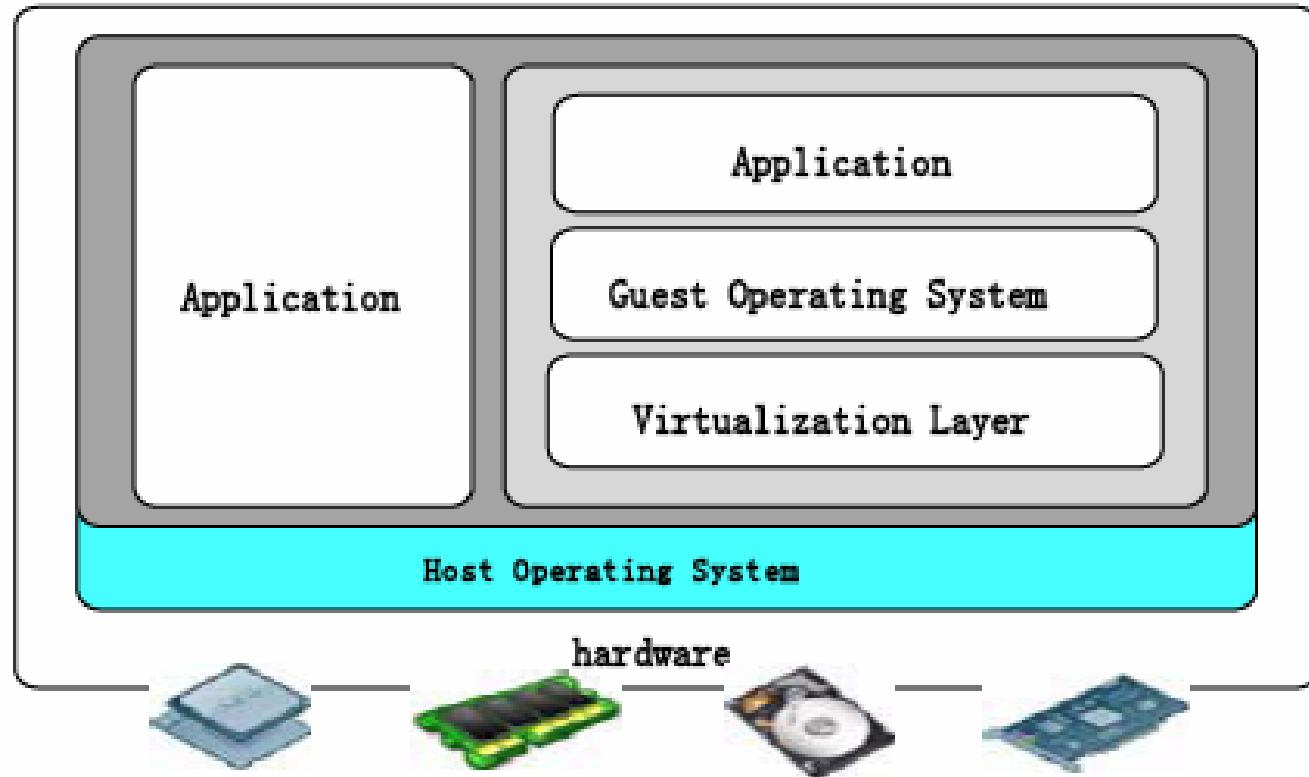


Figure 6.11 A hosted VM installs a guest OS on top of the host OS. This differs from the full virtualization architecture shown in Figure 6.9.

Para-virtualization

In para-virtualization, the guest operating system has to be modified. Para-virtualization provides specially defined ‘hooks’ to undertake some tasks in the host and guest operating systems, which would otherwise have been done in a virtual environment, which is slower.

The VMM in a para-virtualized platform is simpler because the critical tasks are now performed in the operating system rather than by the VMM. Since the virtualization overhead decreases, the performance increases.

Some of the disadvantages are the compatibility and the portability it reduces because of the modified operating system. Also, the cost of maintenance is high because of the deep OS modifications.

Some examples of para-virtualization are KVM and the VMware ESX, and vSphere.

Table 3.5. VMware hardware virtualization (hypervisors) and hosted software for virtualization (<http://vmware.com/products/vsphere/>, 2016)

Category	VMware software products or third party software
Native (Hypervisor)	Adeos, CP/CMS, Hyper-V, KVM (Red Hat Enterprise Virtualization), LDoms / Oracle VM Server for SPARC, LynxSecure, SIMMON, VMware ESXi (VMware vSphere, vCloud), VMware Infrastructure, Xen XenClient), z/VM
Hosted (Specialized)	Basilisk II, bhyve, Bochs, Cooperative Linux, DOSBox, DOSEMU, LLinux, Mac-on-Linux, Mac-on-Mac, SheepShaver, SIMH, Windows on Windows, Virtual DOS machine, Win4Lin
Hosted (Independent)	Microsoft Virtual Server, Parallels Workstation, Parallels Desktop for Mac, Parallels Server for Mac, PearPC, QEMU, VirtualBox, Virtual Iron, VMware Fusion, VMware Player, VMware Server VMware Workstation, Windows Virtual PC
Other tools	Ganeti, oVirt, Virtual Machine Manager

VMWare ESX server for para-virtualization

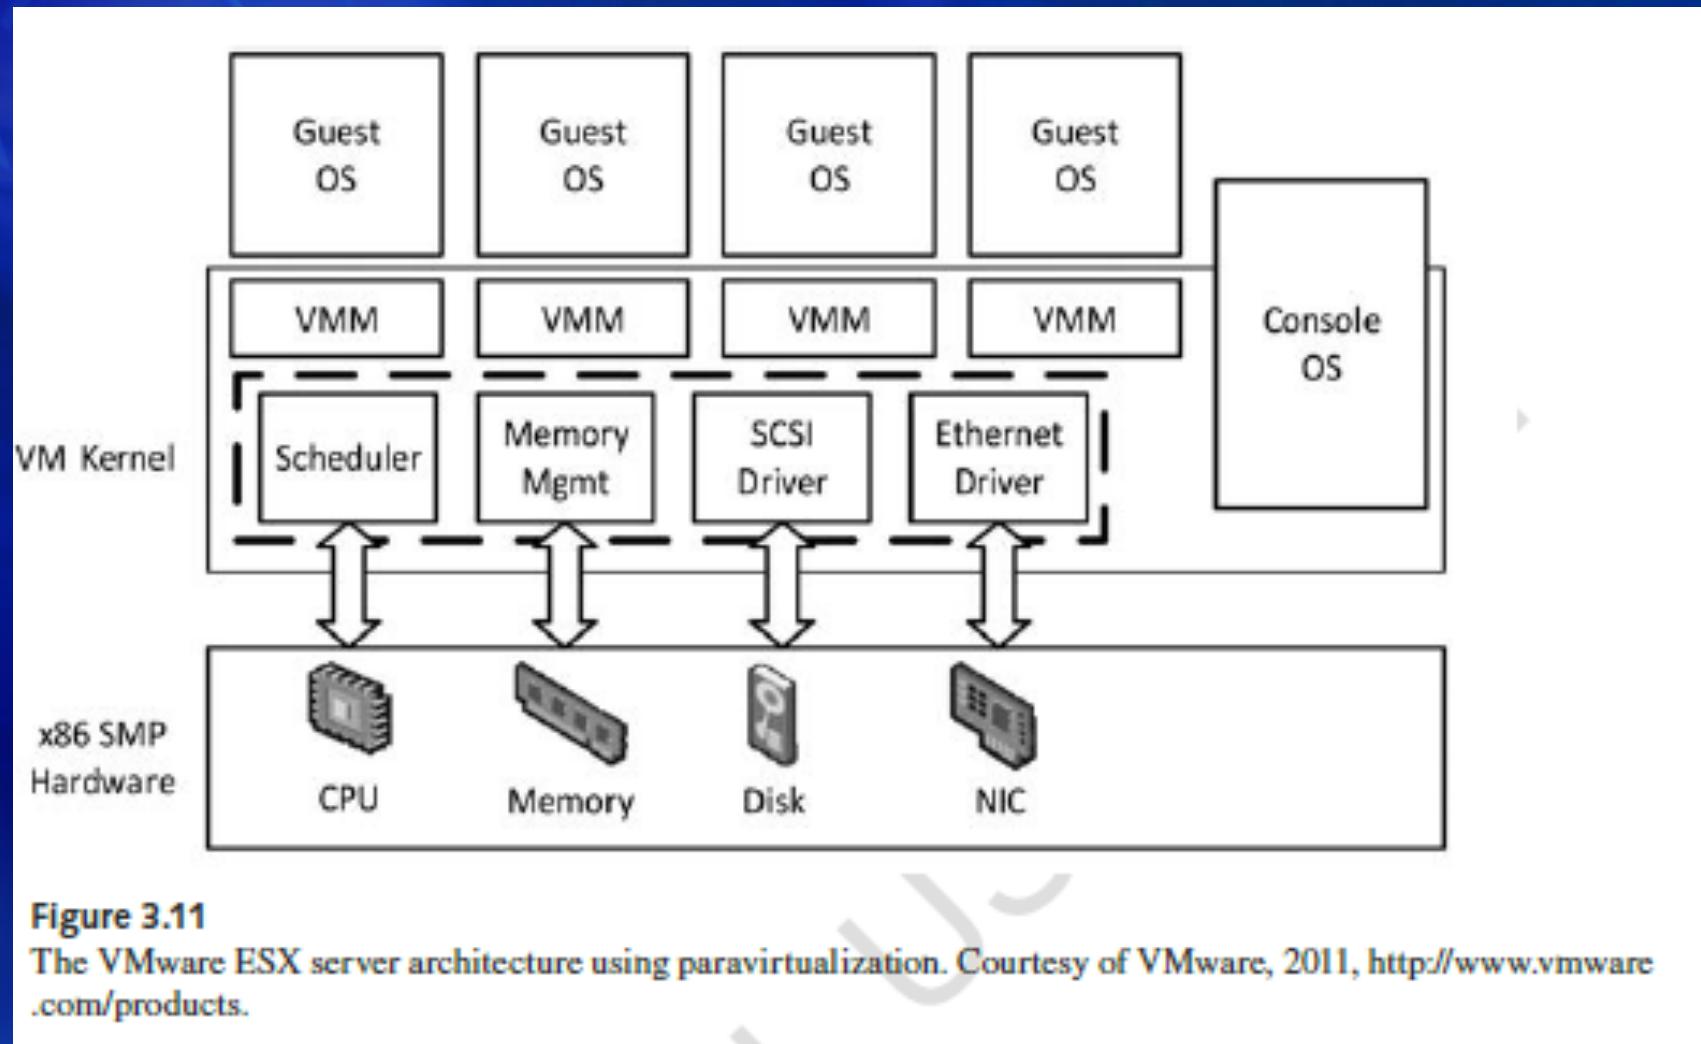


Figure 3.11

The VMware ESX server architecture using paravirtualization. Courtesy of VMware, 2011, <http://www.vmware.com/products>.

HW and SW support for CPU, memory and I/O virtualization

- **CPU virtualization demands hardware-assisted traps of sensitive instructions by the VMM**
- **Memory virtualization demands special hardware support (shadow page tables by VMWare or extended page table by Intel) to help translate virtual address into physical address and machine memory in two stages**
- **I/O virtualization is the most difficult one to realize due to the complexity of I/O service routines and the emulation needed between the guest OS and host OS**

Virtualization support at Intel

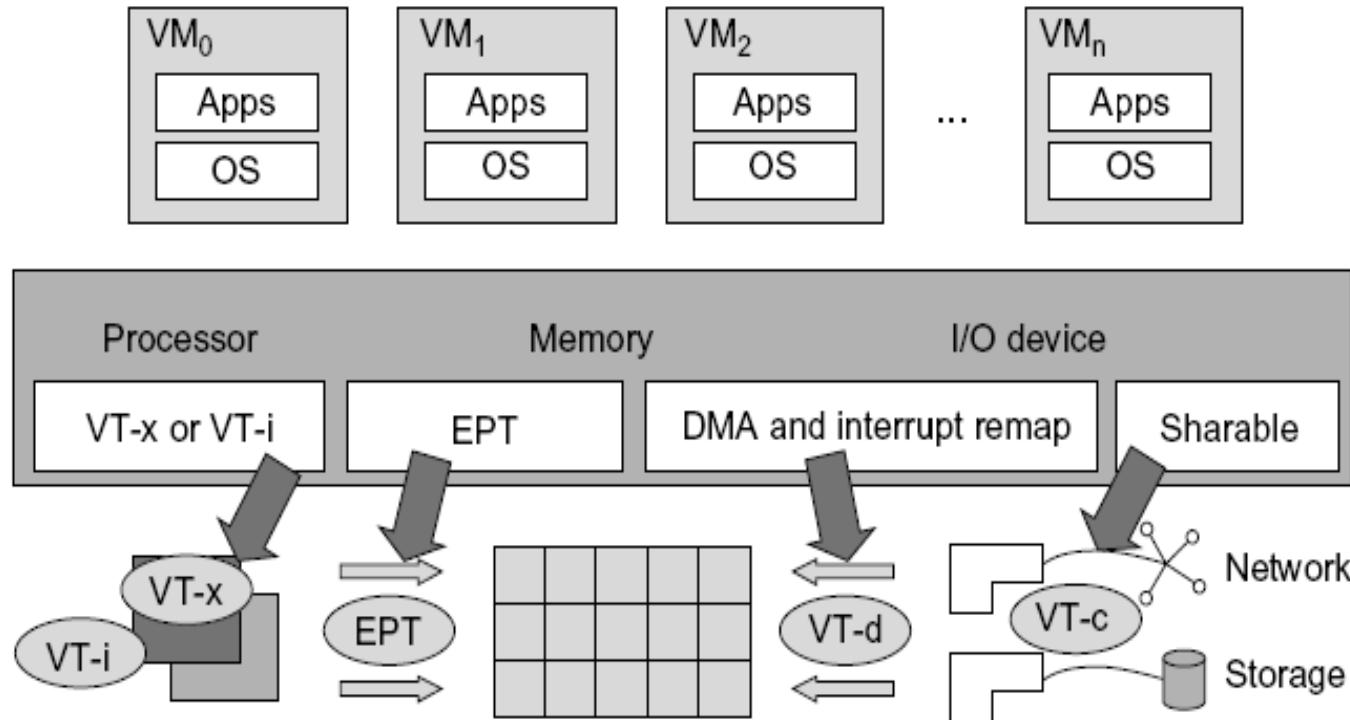


FIGURE 3.10

Intel hardware support for virtualization of processor, memory, and I/O devices.

(Modified from [68], Courtesy of Lihong Chen, USC)

Lecture 6: Docker engine and containers for virtualization at the Linux Kernel Level

- Docker engine works with host OS to produce application software containers
- The App software container does not use a guest OS, thus it is highly scalable
- Virtual clusters are studied with the use either VMs or containers

Virtualization at the operating system level

An abstraction layer between traditional OS and user applications. This virtualization creates isolated containers on a single physical server and the OS-instance to utilize the hardware and software in datacenters. Typical systems: Jail / Virtual Environment / Ensim's VPS / FVM

Advantage: has minimal startup/shutdown cost, low resource requirement, and high scalability; synchronizes VM and host state changes

Shortcoming and limitation: all VMs at the operating system level must have the same kind of guest OS; poor application flexibility and isolation

Virtualization at OS level

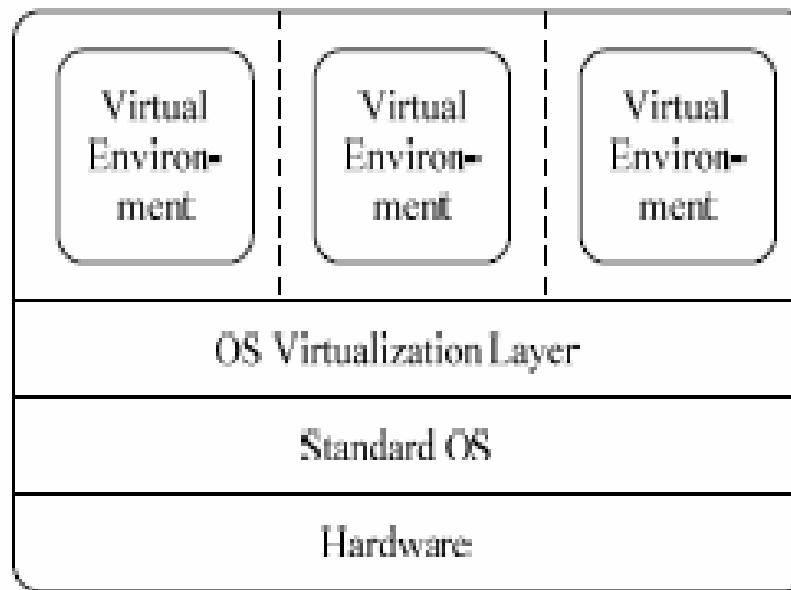


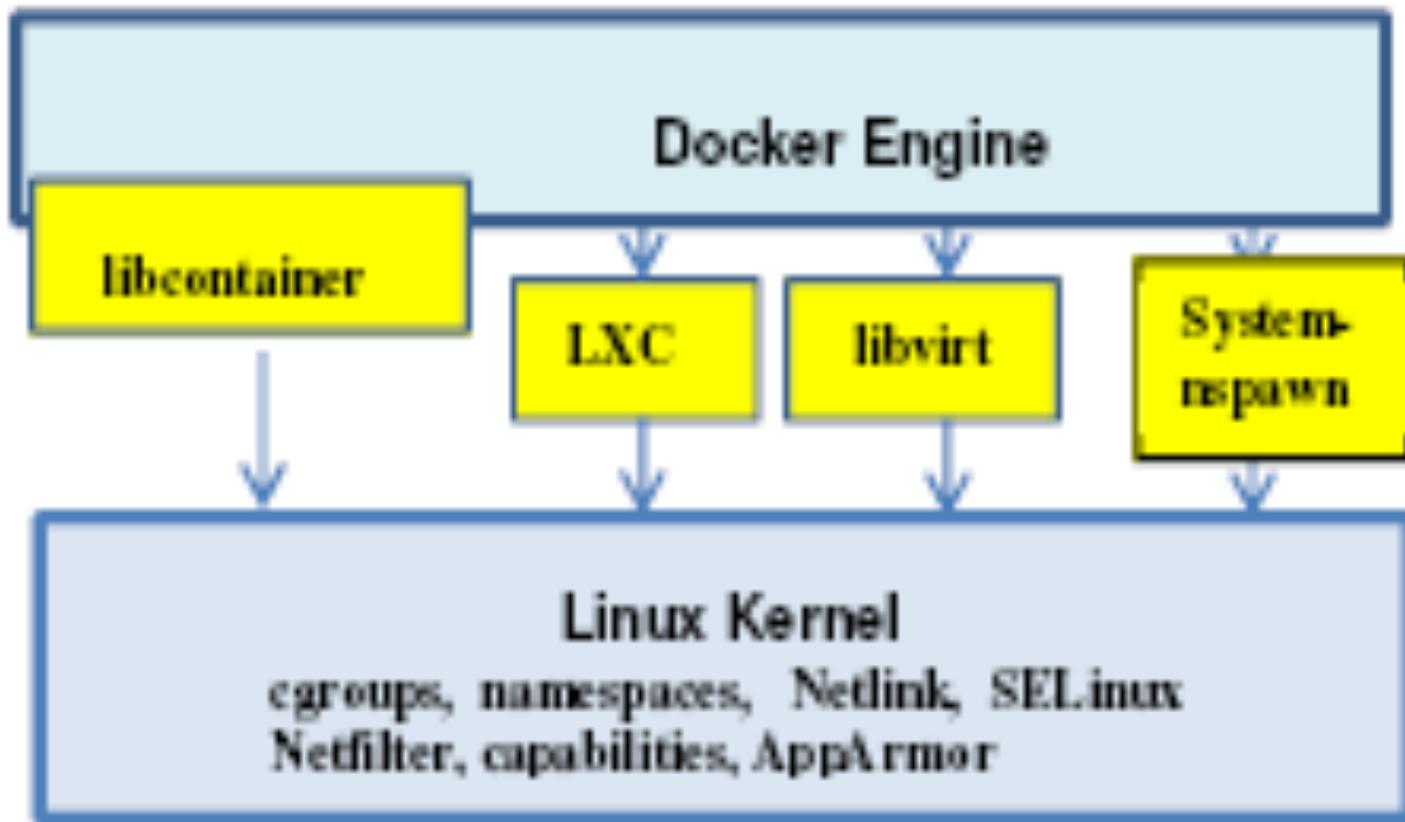
Figure 6.3 The virtualization layer is inserted inside an OS to partition the hardware resources for multiple VMs to run their applications in virtual environments

Advantages of OS extension for virtualization

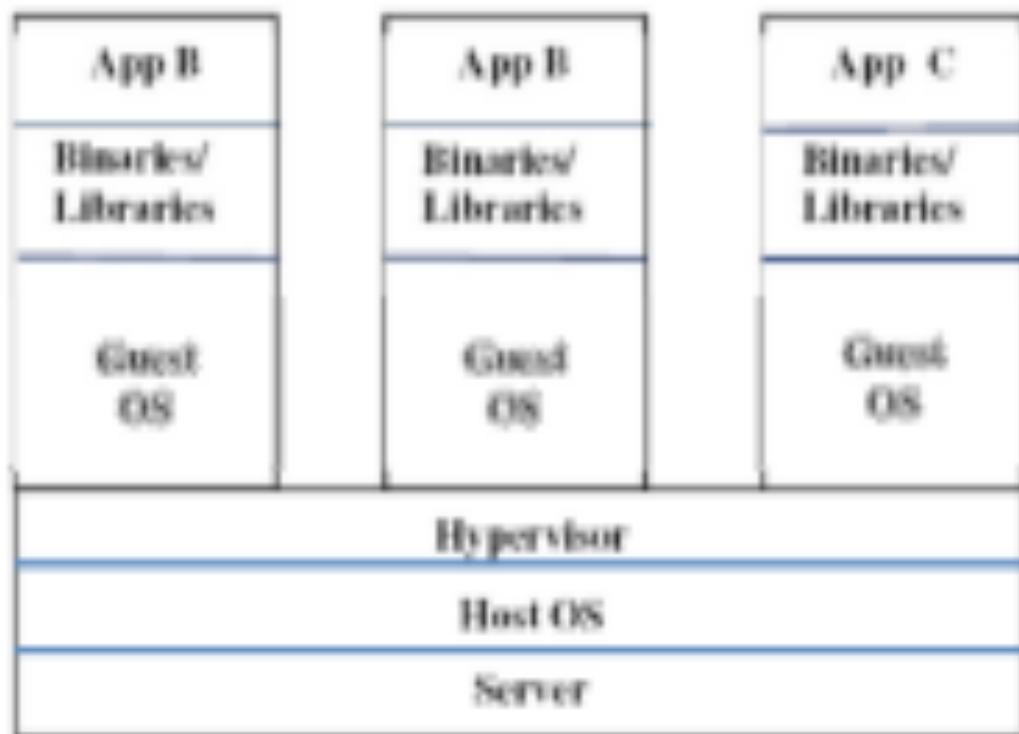
- 1. VMs at OS level has minimum startup/shutdown costs**
- 2. OS-level VM can easily synchronize with its environment**

Disadvantage of OS Extension for Virtualization

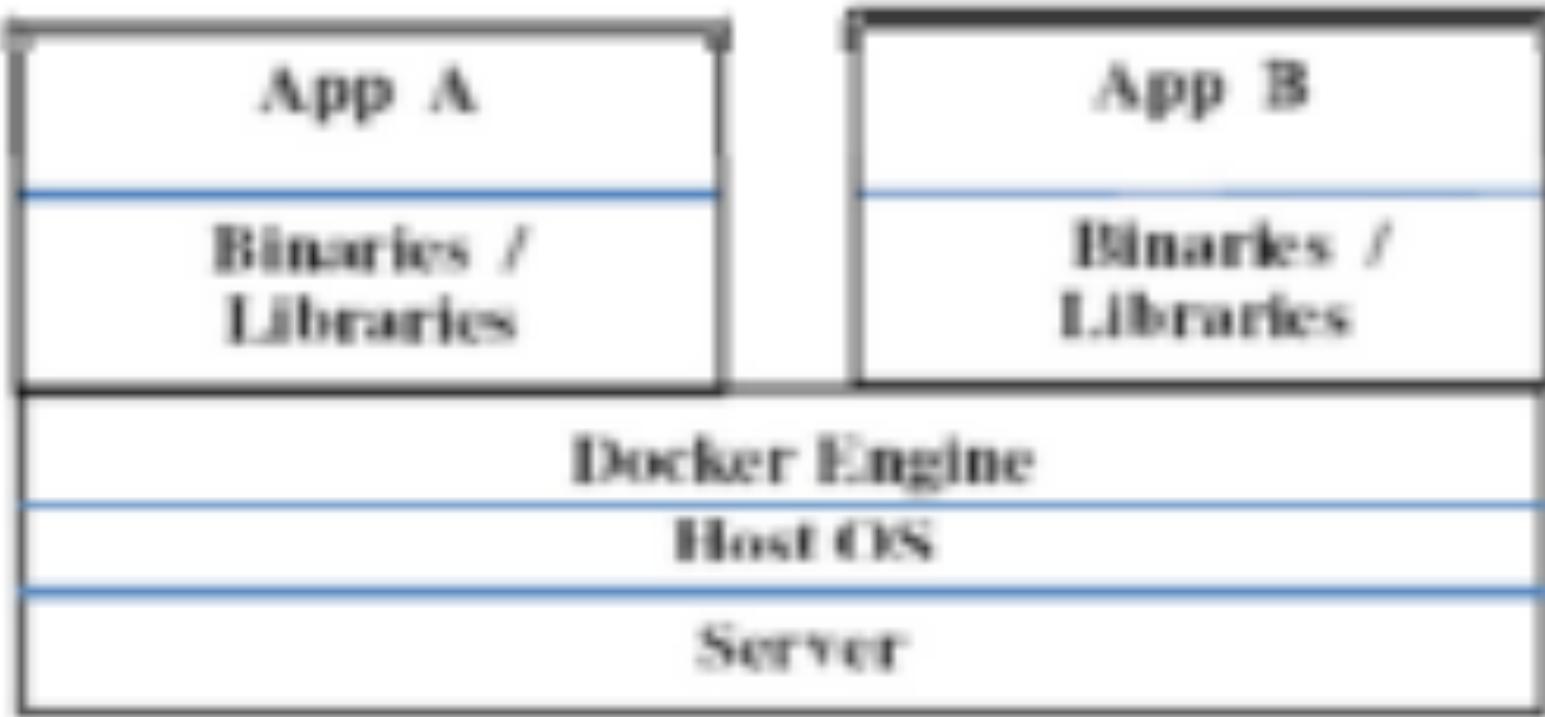
All VMs in the same OS container must have the same or similar guest OS, which restricts application flexibility of different VMs on the same physical machine



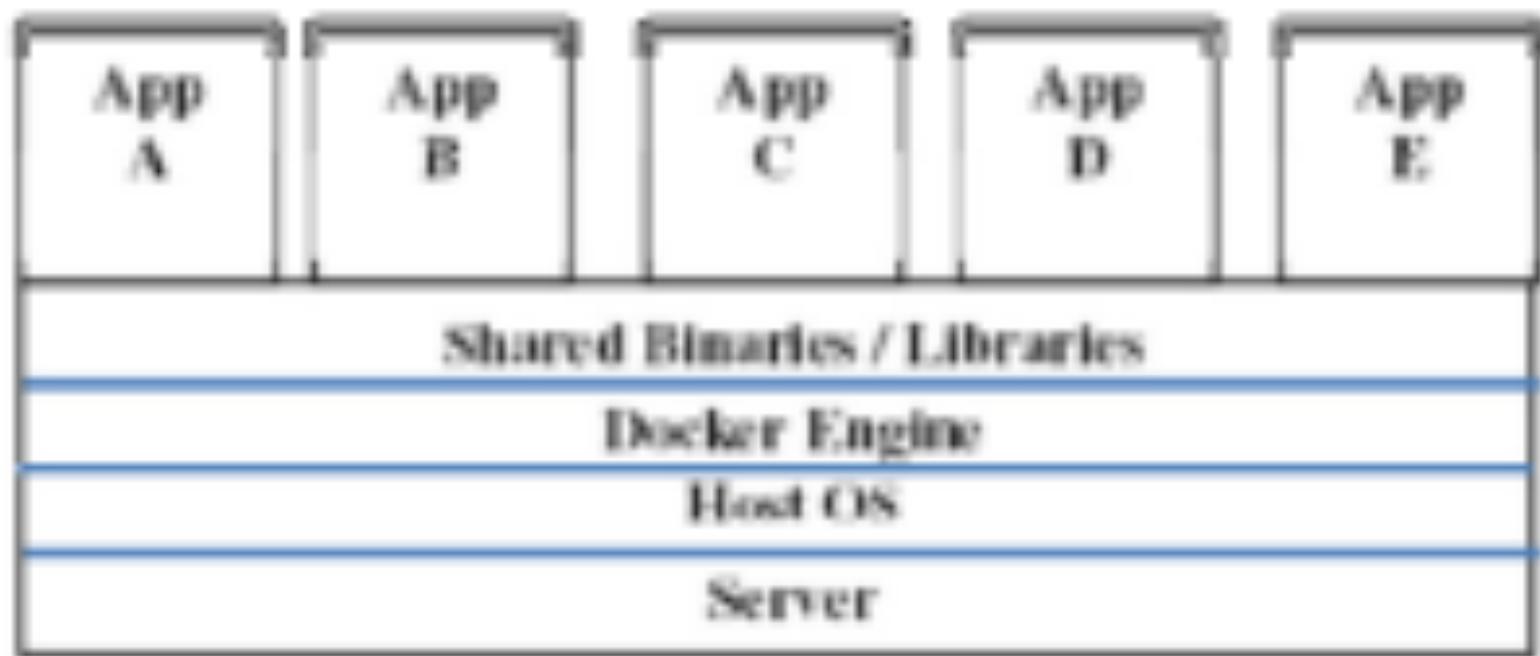
The Docker engine access the Linux kernel features for isolated virtualization of different application containers.



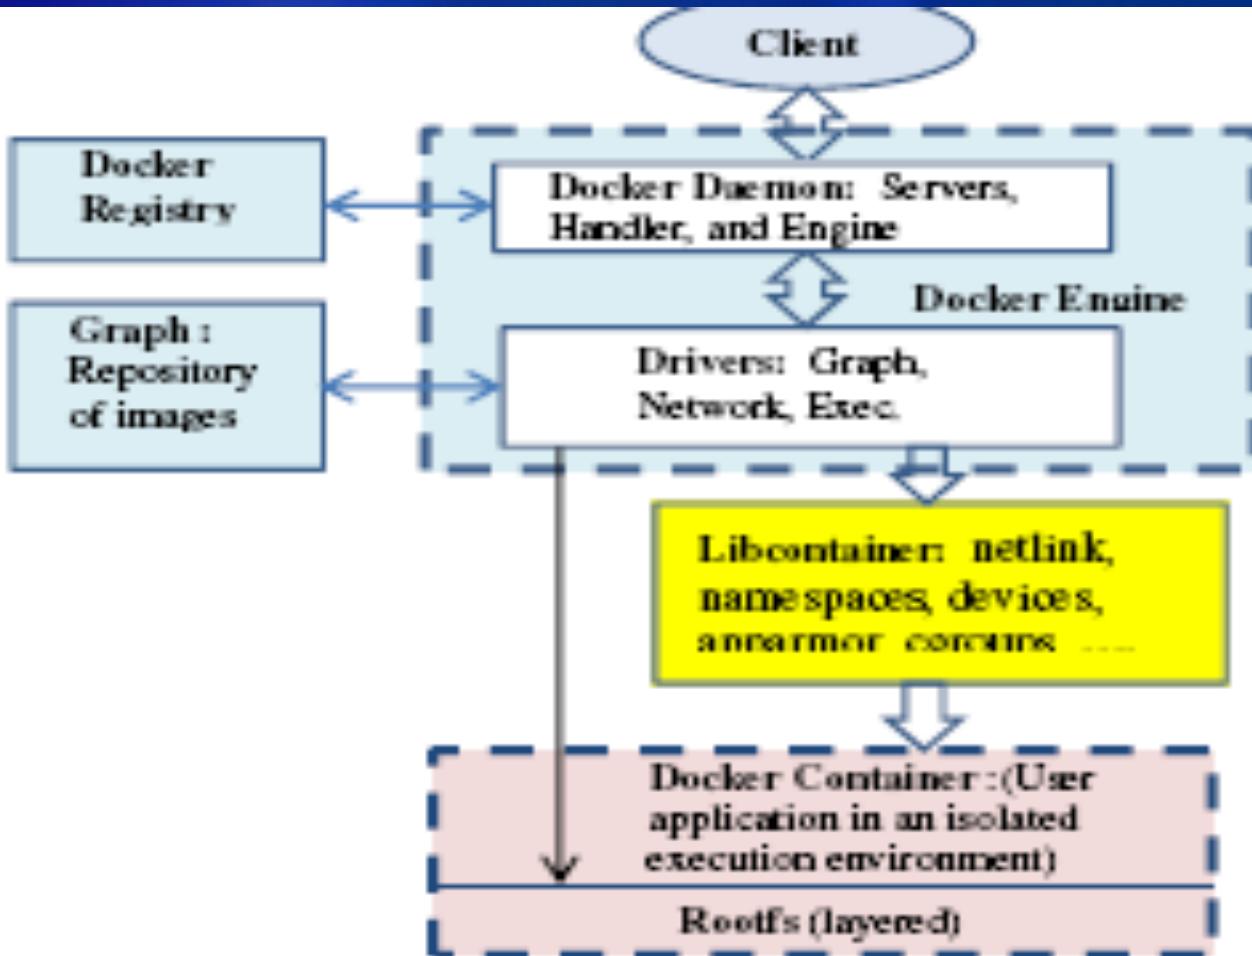
(a). Three hypervisor-created virtual machines (VMs) on the same hardware host, each VM is heavily loaded with its own guest OS and specific binaries and libraries.



(b). Each container is loaded with its own binaries/libraries which are not shared.



- (c) The Docker engine creates many light-weight app containers, which are isolated , but share OS and ,



Docker schematic block diagram illustrating the process of creating and management of Docker containers for isolated execution, separately.

Table 1.9: Comparison Between Hypervisor-created VMs and Docker Containers

VM Type	Strength and Weakness	Suitable Applications
Hypervisor-created Virtual Machines	Higher app flexibility in launching different OS apps, but demand more memory and overhead to create and launch the VMs	More suitable for use in multiple apps without orchestration. VMs appeal to run on wide variety of operating systems.
Docker Containers	Light-weight application containers with low overhead to create and running with better protection under an isolated execution environment.	More suitable for scalable use of the same app in multiple copies under orchestration. Works better with a particular OS version. It may save the operation costs on clouds.

Assessing the use of virtual machines and docker containers in today's clouds

- Both VMs and software containers will co-exist for some time in today's clouds
- The VMs have high software portability on different types of hardware platforms
- VMs are heavily weighted with the use of heavy duty guest OS. This may weaken its acceptance in the future
- The docker containers are light-weight and more cost-effective to implement and to apply in scalable applications. Eventually, most clouds will use containers over Linux hosts

Software tools for container management

Table 3.11 Host Provisioning and Container Scheduling Tools

Tool Name	Brief Description of Tool Functionality
fleet	Scheduling and cluster management component of CoreOS
marathon	Scheduling and service management component a Mesosphere installation
Swarm	Docker's robust scheduler to spin up containers on provisioned hosts
mesos	Apache mesos abstracts and manage the resources of all hosts in a cluster
Kubernetes	Google's scheduler over the containers running on your cloud infrastructure
compose	Docker's tool that allows group management of containers, declaratively

Virtualization at library support level

It creates execution environments for running alien programs on a platform rather than creating VM to run the entire operating system. It is done by API call interception and remapping. Typical systems: Wine, WAB, LxRun, VisualMainWin

Advantage: It has very low implementation effort

Shortcoming & limitation: poor application flexibility and isolation

Virtualization with middleware/library support

Table 3.4 Middleware and Library Support for Virtualization

Middleware or Runtime Library and References or Web Link	Brief Introduction and Application Platforms
WABI (http://docs.sun.com/app/docs/doc/802-6306)	Middleware that converts Windows system calls running on x86 PCs to Solaris system calls running on SPARC workstations
Lxrun (Linux Run) (http://www.ugcs.caltech.edu/~steven/lxrun/)	A system call emulator that enables Linux applications written for x86 hosts to run on UNIX systems such as the SCO OpenServer
WINE (http://www.winehq.org/)	A library support system for virtualizing x86 processors to run Windows applications under Linux, FreeBSD, and Solaris
Visual MainWin (http://www.mainsoft.com/)	A compiler support system to develop Windows applications using Visual Studio to run on Solaris, Linux, and AIX hosts
vCUDA (Example 3.2) (IEEE IPDPS 2009 [57])	Virtualization support for using general-purpose GPUs to run data-intensive applications under a special guest OS

User-application level

It virtualizes an application as a virtual machine. This layer sits as an application program on top of an operating system and exports an abstraction of a VM that can run programs written and compiled to a particular abstract machine definition. Typical systems: JVM , .NET CLI, Panot

Advantage: best application isolation

Shortcoming & limitation: low performance, low application flexibility and high implementation complexity.

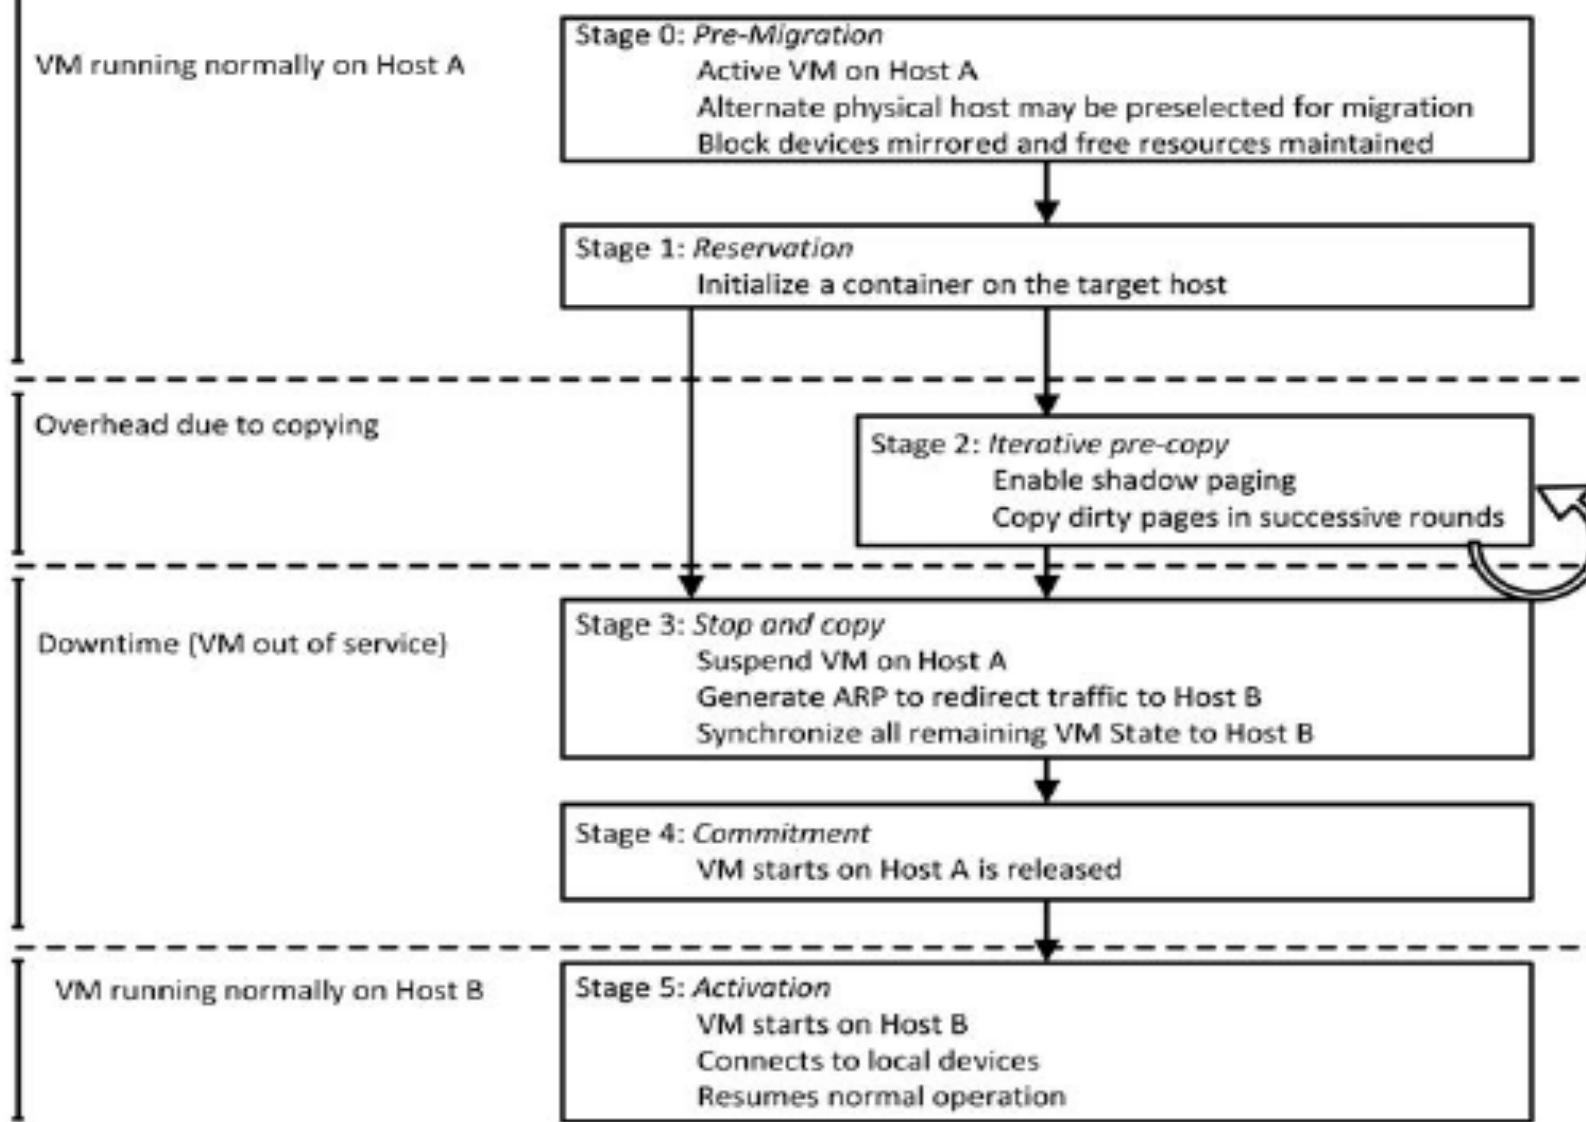


Figure 3.18

Steps for live migration of VMs from host computer A to another host B. Reprinted with permission from C. Clark et al., "Live Migration of Virtual Machines," *Proc. of the Second Symposium on Networked System Design and Implementation* (NSDI'05), 2005.

Live migration of virtual machines

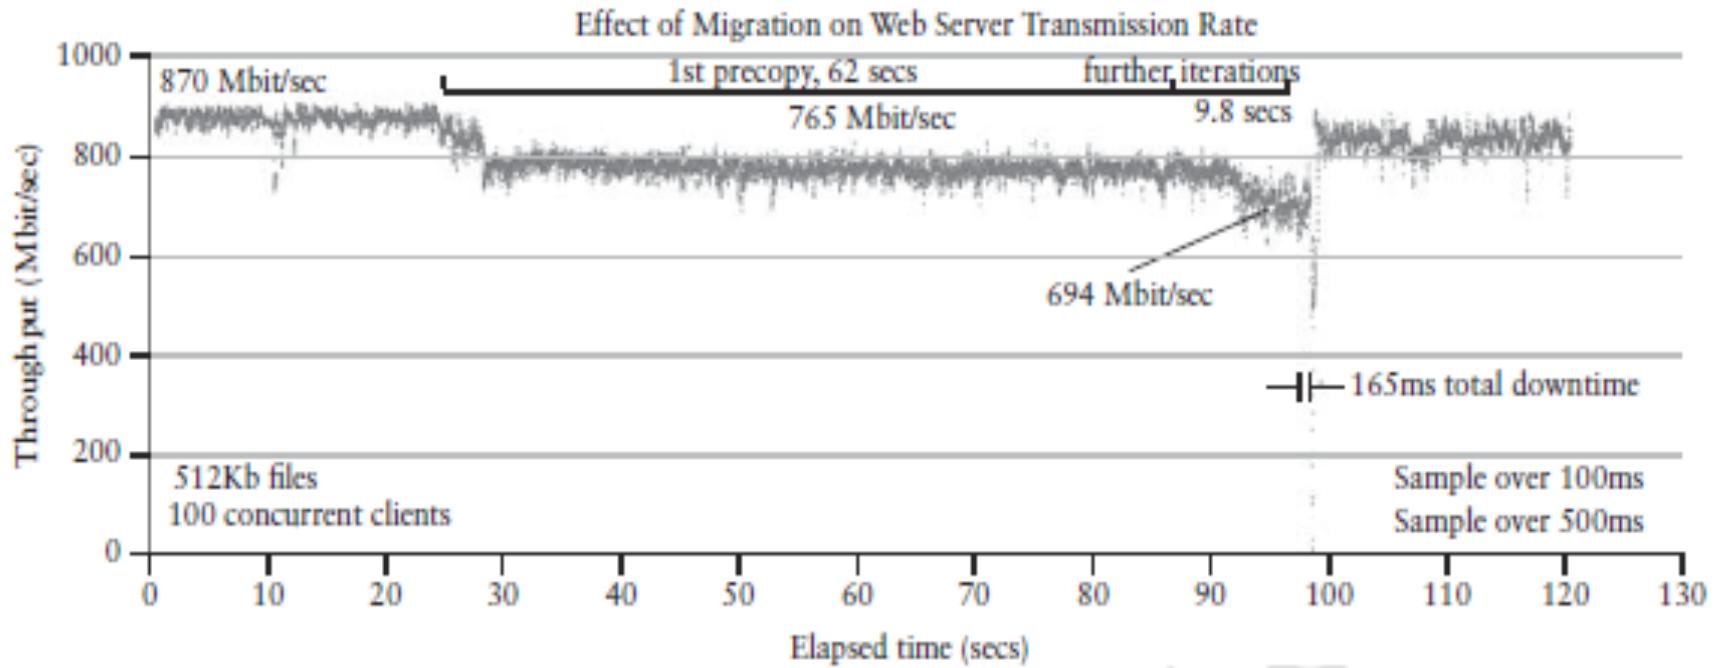


Figure 3.19

Effect on data transmission rate of a VM migrated from one failing web server to another. The downtime of 165 ms was observed. Reprinted with permission from C. Clark et al., "Live Migration of Virtual Machines," *Proc. of the Second Symposium on Networked Systems Design and Implementation* (NSDI'05), 2005.

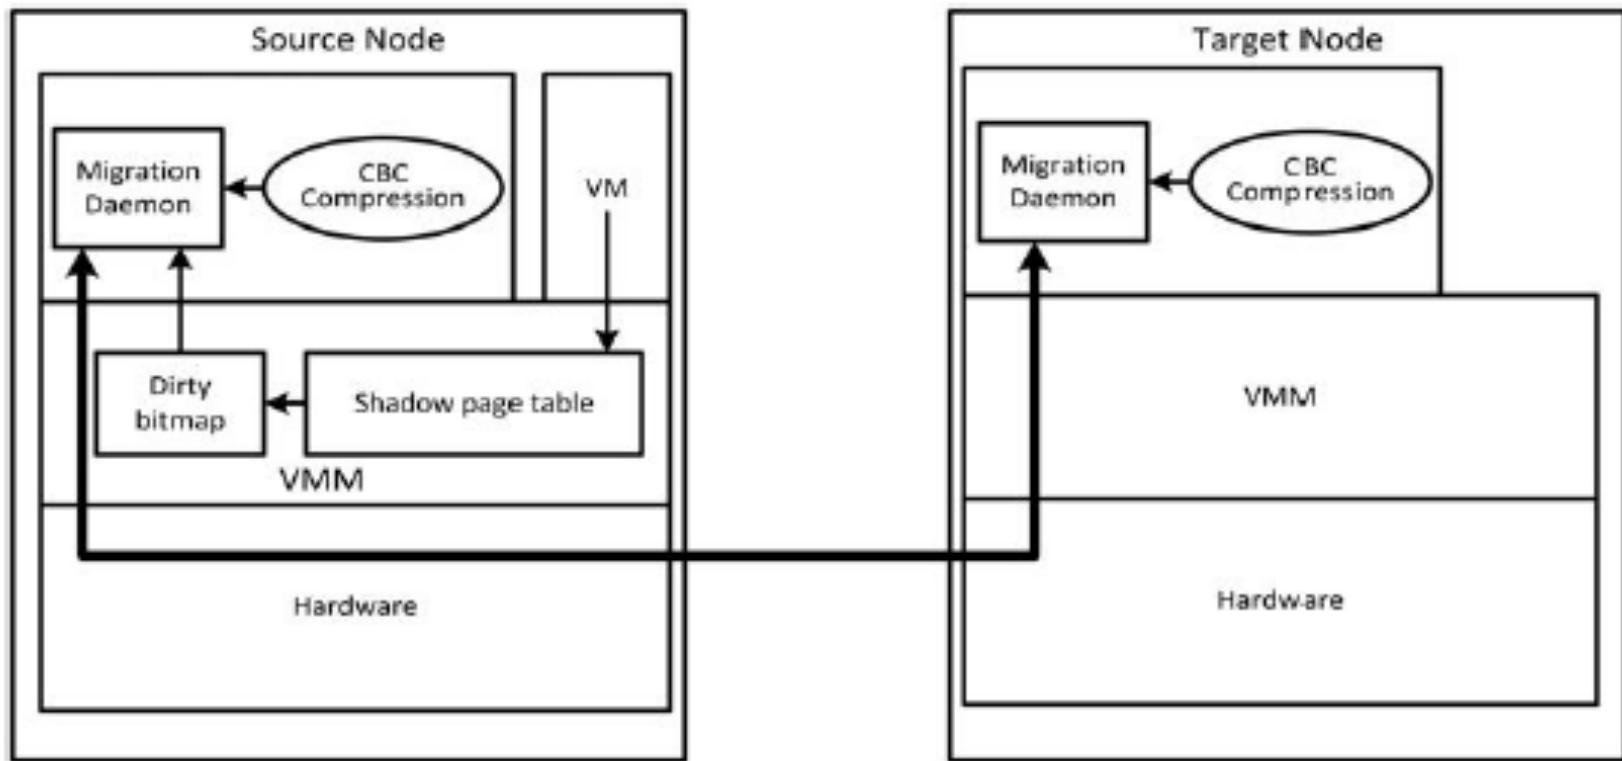


Figure 3.20

Live migration of VM from the Domain0 to that of an Xen-enabled target host.

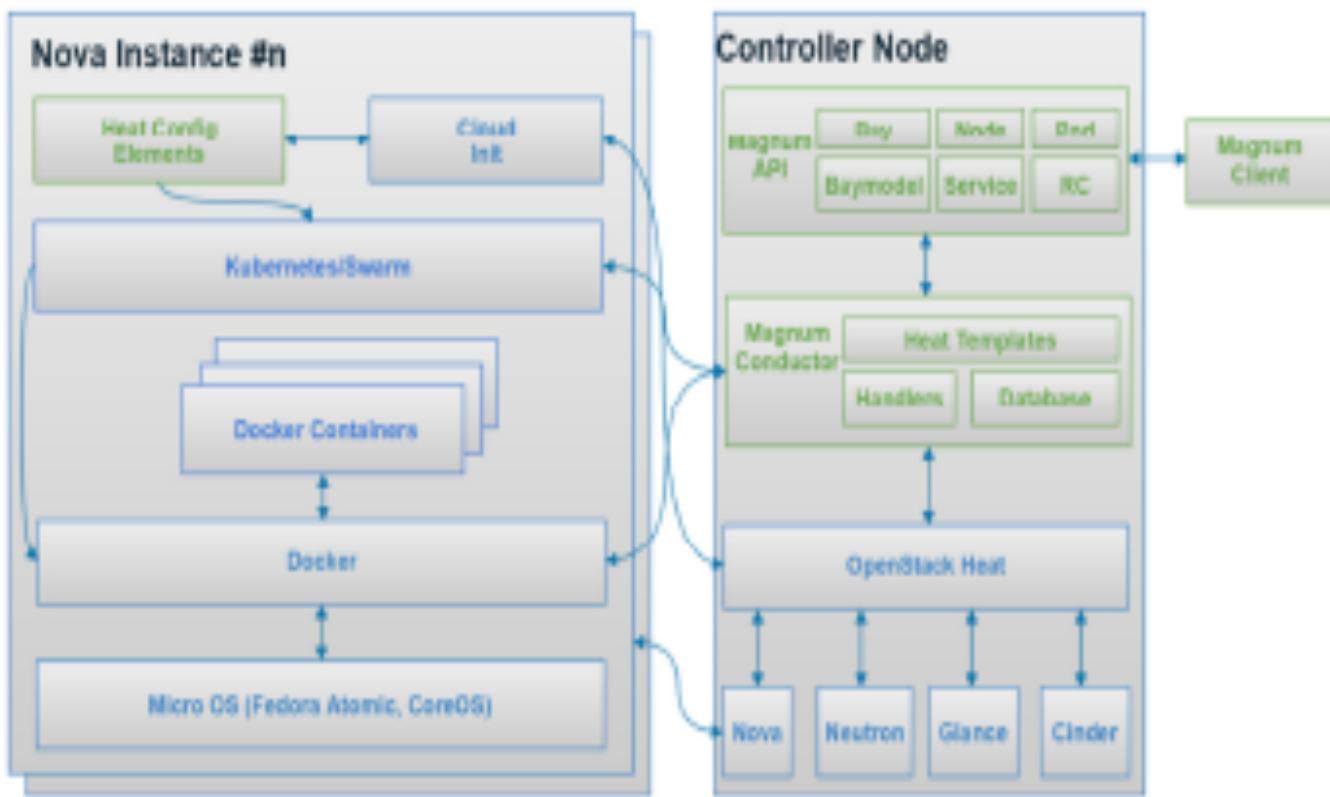
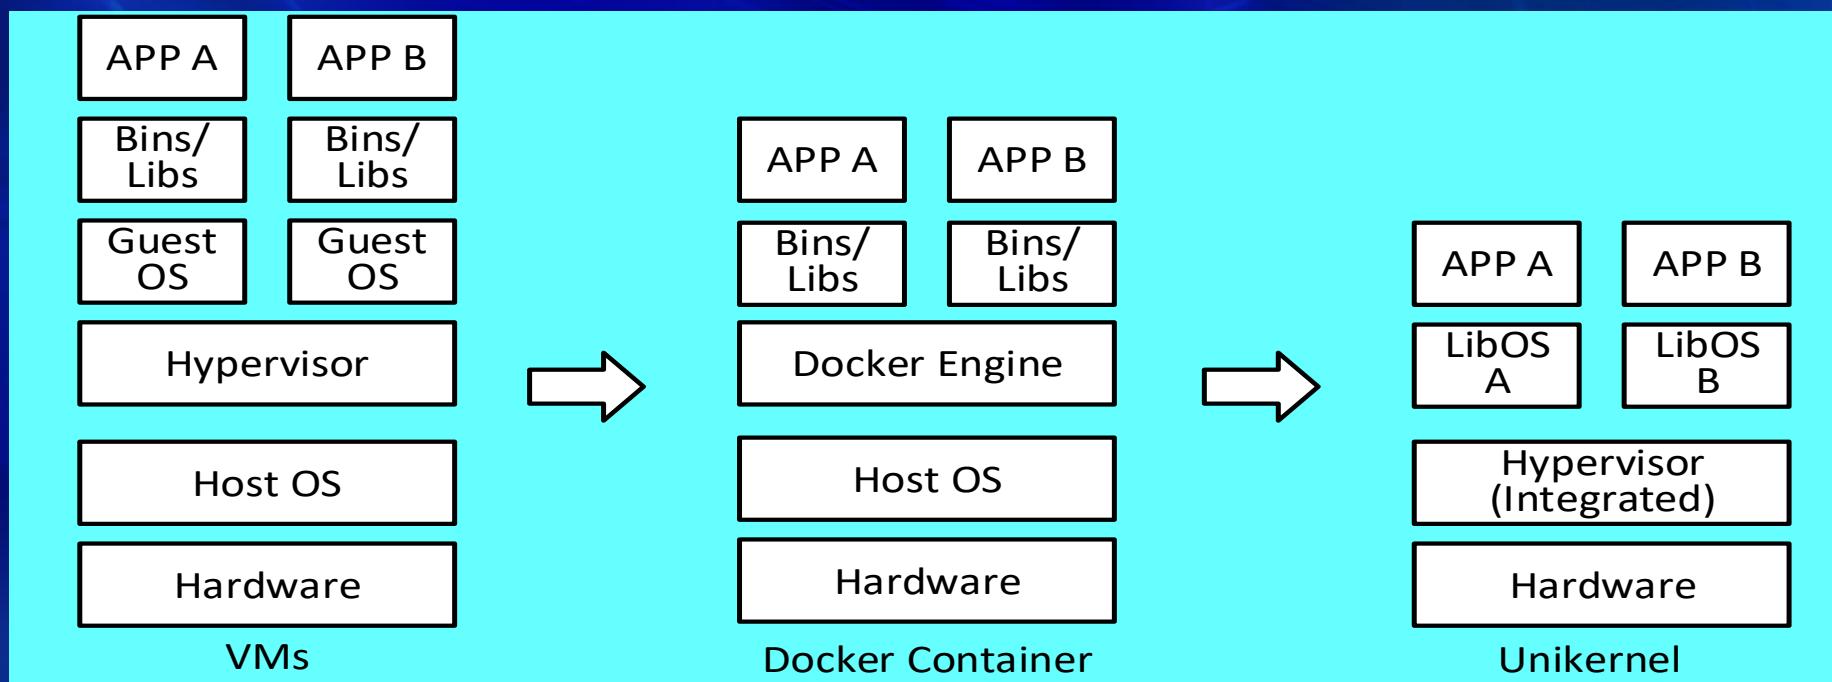


Figure 3.22 Docker container orchestration illustrated by using OpenStack Magnum to deploy container clusters in multiple Nova instances (Courtesy of <https://wiki.openstack.org/wiki/Magnum>, retrieved Aug. 2015).

Figure 3.17. Architectural evolution from VMs to docker containers and unikernels



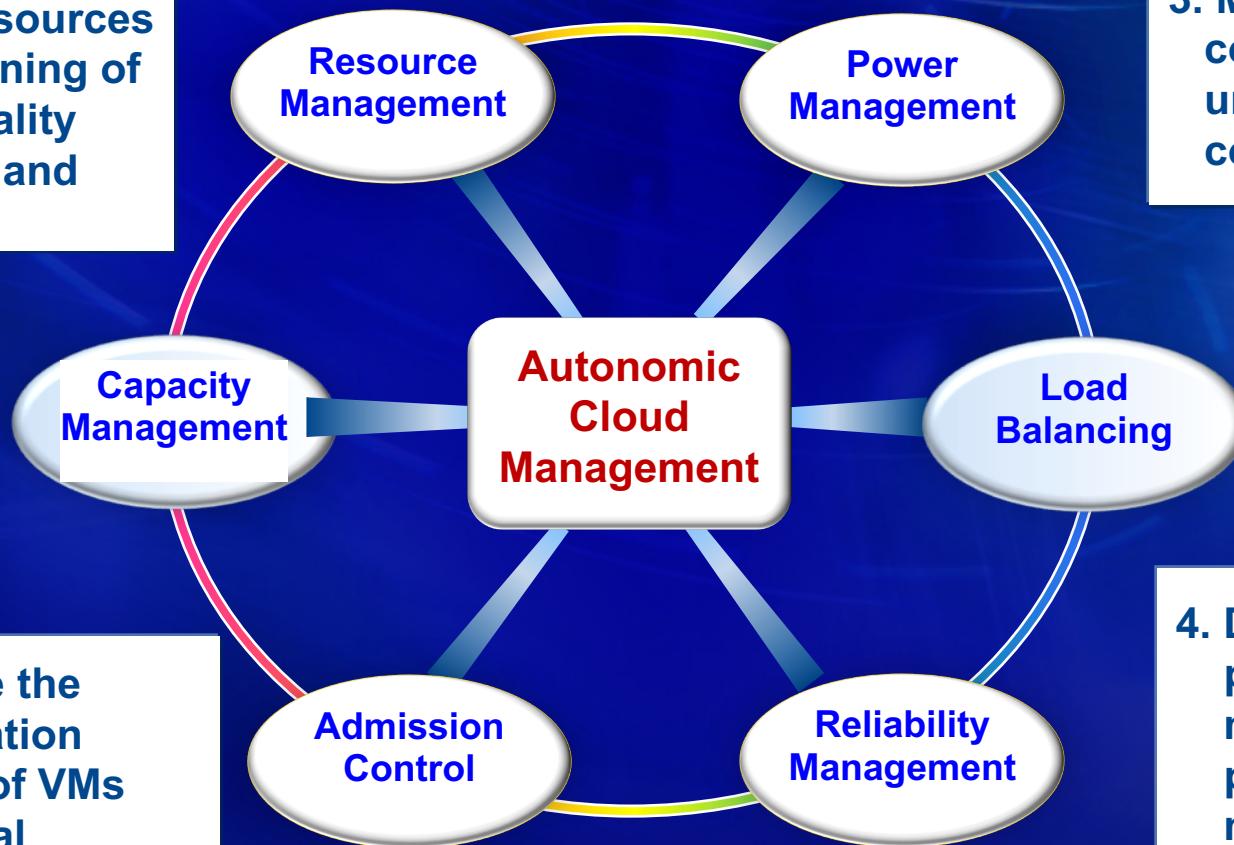
Switching from x86 to ARM, Power, and Sparc as building nodes in clouds

- The x86 processors are used in almost all server clusters deployed in cloud platforms. There is a rise of using ARM, Power and SPARC processors in virtualized cloud computing
- In 2013, IBM announced one billion investment in upgrading its Power series for Linux applications. IBM has provided support plans for commercial clouds to switch to Linux-based Power hosts
- In 2015, Oracle started to shift towards cloud computing. ARM processors used in many smartphone have demonstrated their strength in low power consumption and low cost to deliver high performance
- Dell, HP, Microsoft, and Amazon are all developing ARM servers cloud constructions. Future clouds may enter an era of keen competition among x86, ARM, Power and SPARC processors

Autonomic Cloud Management

Develop methodologies and tools to automate the process of cloud management in 4 objectives

1. Manage resources to provisioning of service quality assurance and adaptation



3. Manage energy consumption under SLA constraints

2. Automate the configuration process of VMs and virtual clusters

4. Develop fault prediction models for proactive failure management

Virtualization for datacenter automation to serve millions of clients, simultaneously

- **Server consolidation in virtualized datacenter**
- **Virtual storage provisioning and deprovisioning**
- **Cloud operating systems for virtual datacenters**
- **Trust management in virtualized datacenters**

Table 3.11

Open-source software for cloud computing (except the vSphere 6)

Software	Cloud Type, License	Language Used	Linux/ Windows	EC2/S3 Compatibility	XEN/KVM/ VMWare
Eucalyptus	IaaS, Rackspace	Java, C	Yes/Yes	Yes/Yes	Yes/Yes/Yes
Nimbus	IaaS, Apache	Java, Python	Unknown	Yes/No	Yes/Yes/ Unknown
Cloud Foundry	PaaS, Apache	Ruby, C	Yes/No	Yes/No	Yes/Yes/Yes
OpenStack	IaaS, Apache	Python	Yes/Unknown	Yes/Yes	Yes/Yes/ Unknown
OpenNebula	IaaS, Apache	C, C++, Ruby, Java, Lex, YaaS, Shell Script	Yes/Unknown	Yes/Unknown	Yes/Yes/ Unknown
AppScale	Unknown	Unknown	Unknown	Yes/Yes	Yes/Yes/Yes
vSphere 6	Unknown	Unknown	Yes/Yes	Yes/Yes	Yes/Yes/Yes

In Prob. 3.17, you are supposed to dig out more details from VMWare published material. Visit their website.

Cloud software packages and features

Software	Cloud Type	License(s)	Language	Linux/ Windows	EC2/S3	Xen/KVM/ VMWare	Virtual Box	OCCI/ vCloud
Fluid Operations	IaaS, Paas, LaaS, SaaS, TaaS, Daas, BaaS	Proprietary	Java, C	Yes/Yes	Yes/No	Yes/Yes/ yes	?	No /Yes
AppIScale	Paas	BSD	Python, Ruby, Go	? / ?	Yes/ Yes	Yes/Yes/ yes	Yes	? / ?
Cloud Foundry	PaaS	Apache	Ruby, C	Yes/No	Yes/No	Yes/Yes/ yes	Yes	No /Yes
Cloud.com	IaaS	Proprietary, GPLv3	Java, C	? / ?	? / ?	Yes/Yes/ yes	?	? / ?
Eucalyptus	IaaS	Proprietary, GPLv3	Java, C	Yes/No	Yes/ Yes	Yes/Yes/ yes	?	? / ?
Nimbus	IaaS	Apache	Java, Python	? / ?	Yes/No	Yes/ Yes/?	?	? / ?
OpenNebula	IaaS	Apache	C++,C,Ruby, Java, lex, yacc, Shellscript	Yes/ ?	Yes/ ?	Yes/ Yes/?	?	Yes/Yes
OpenStack	IaaS	Apache	Python	Yes/ ?	Yes/ Yes	Yes/ Yes/?	?	? / ?

Source: http://en.wikipedia.org/wiki/Cloud_computing_comparison (read 02/02/2012)

Cloud OS for building private clouds

Table 3.6 VI Managers and Operating Systems for Virtualizing Data Centers [9]

Manager/ OS, Platforms, License	Resources Being Virtualized, Web Link	Client API, Language	Hypervisors Used	Public Cloud Interface	Special Features
Nimbus Linux, Apache v2	VM creation, virtual cluster, www.nimbusproject.org/	EC2 WS, WSRF, CLI	Xen, KVM	EC2	Virtual networks
Eucalyptus Linux, BSD	Virtual networking (Example 3.12 and [41]), www.eucalyptus.com/	EC2 WS, CLI	Xen, KVM	EC2	Virtual networks
OpenNebula Linux, Apache v2	Management of VM, host, virtual network, and scheduling tools, www.opennebula.org/	XML-RPC, CLI, Java	Xen, KVM	EC2, Elastic Host	Virtual networks, dynamic provisioning
vSphere 4 Linux, Windows, proprietary	Virtualizing OS for data centers (Example 3.13), www.vmware.com/ products/vsphere/ [66]	CLI, GUI, Portal, WS	VMware ESX, ESXi	VMware vCloud partners	Data protection, vStorage, VMFS, DRM, HA

Eucalyptus: An open-source cloud operating system

- A software platform developed by Eucalyptus Systems, Inc., (started 2008 and stable release 2010)
- Written in Java, C, running with Linux, can host Linux and Windows VMs
- Uses hypervisors (Xen, KVM and VMWare) and compatible with EC2 and S3 services
- Eucalyptus stands for “Elastic Utility Computing Architecture for Linking Your Programs To Useful Systems”
- For use in developing IaaS-style private cloud or hybrid cloud on computer cluster, working with AWS API
- License: Proprietary or GPLv3 for open-core enterprise edition. Open-source edition available
- Website: www.eucalyptus.com

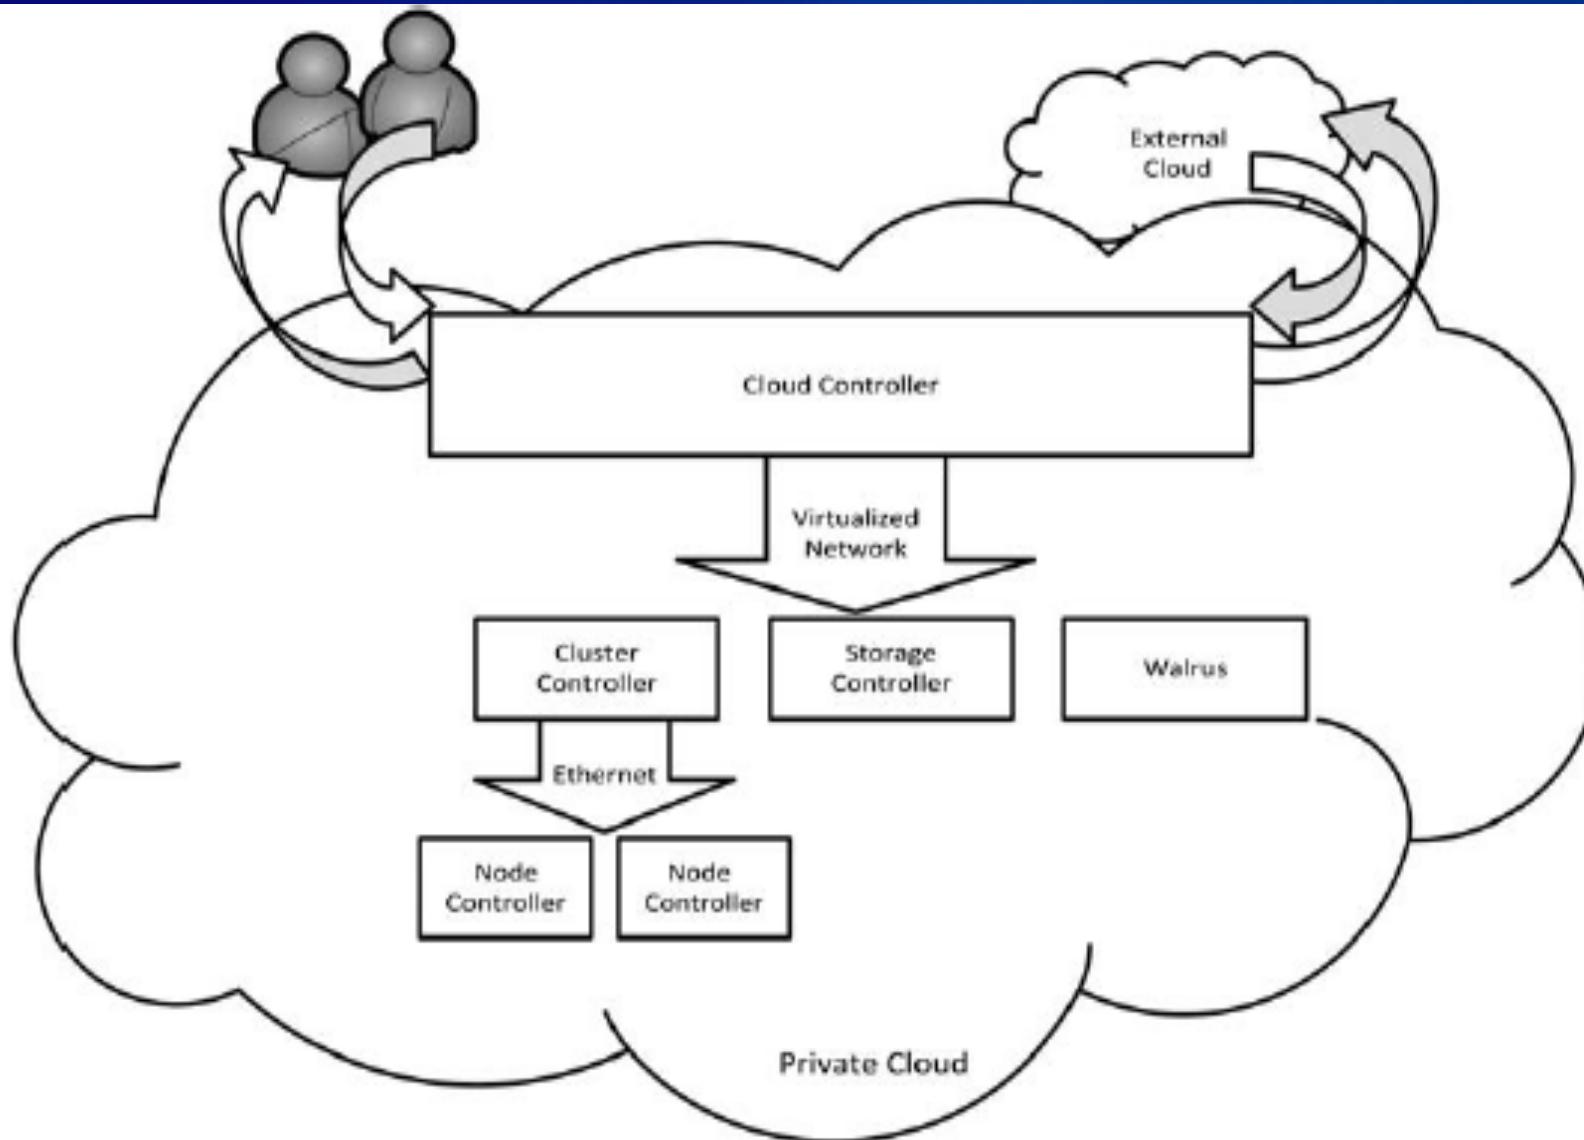


Figure 3.22

The Eucalyptus for building private cloud by establishing virtual network over the VMs linking through Ethernet and the Internet.

OpenStack

main services and components

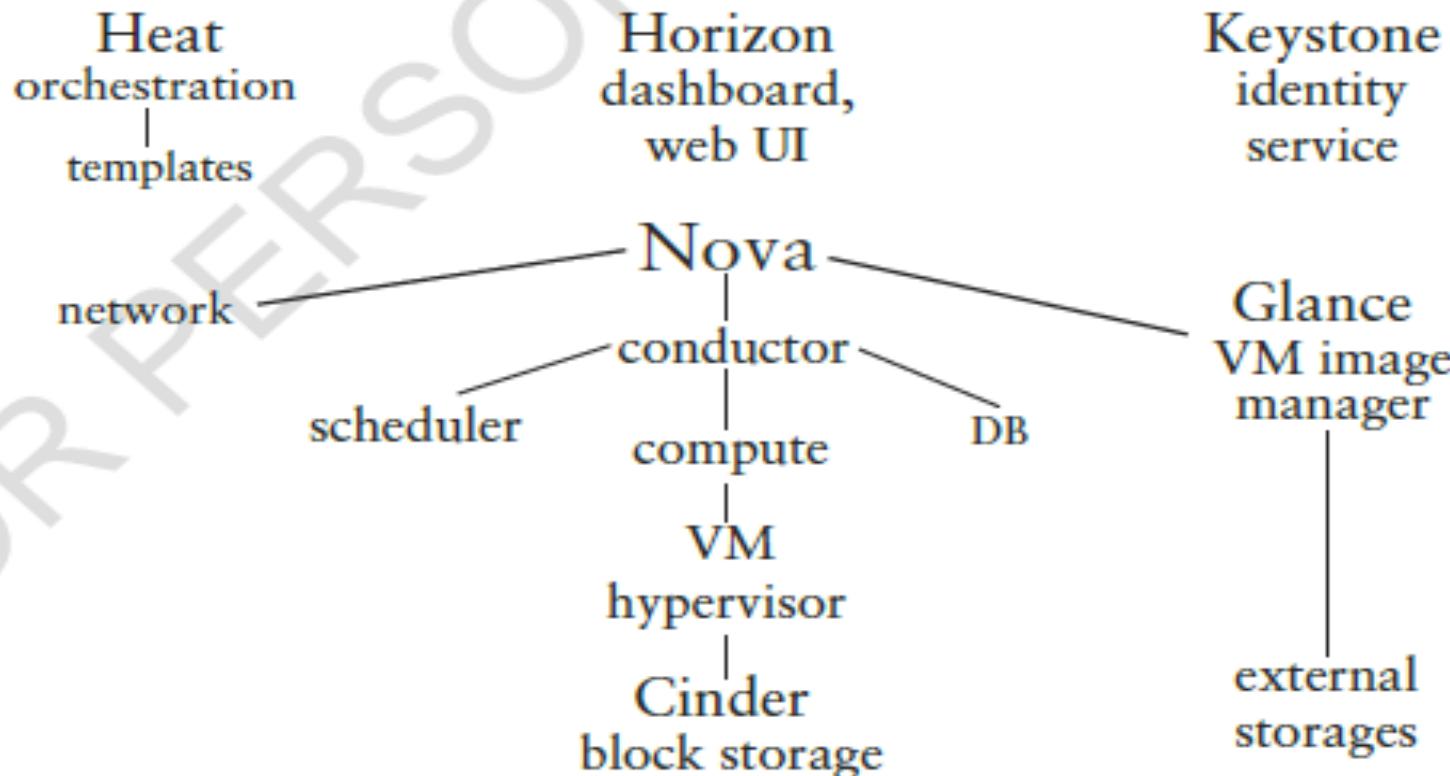


Figure 3.24

OpenStack for constructing private or public clouds in IaaS services. Courtesy of OpenStack, <http://openstack.org>, Apache License 2.0.

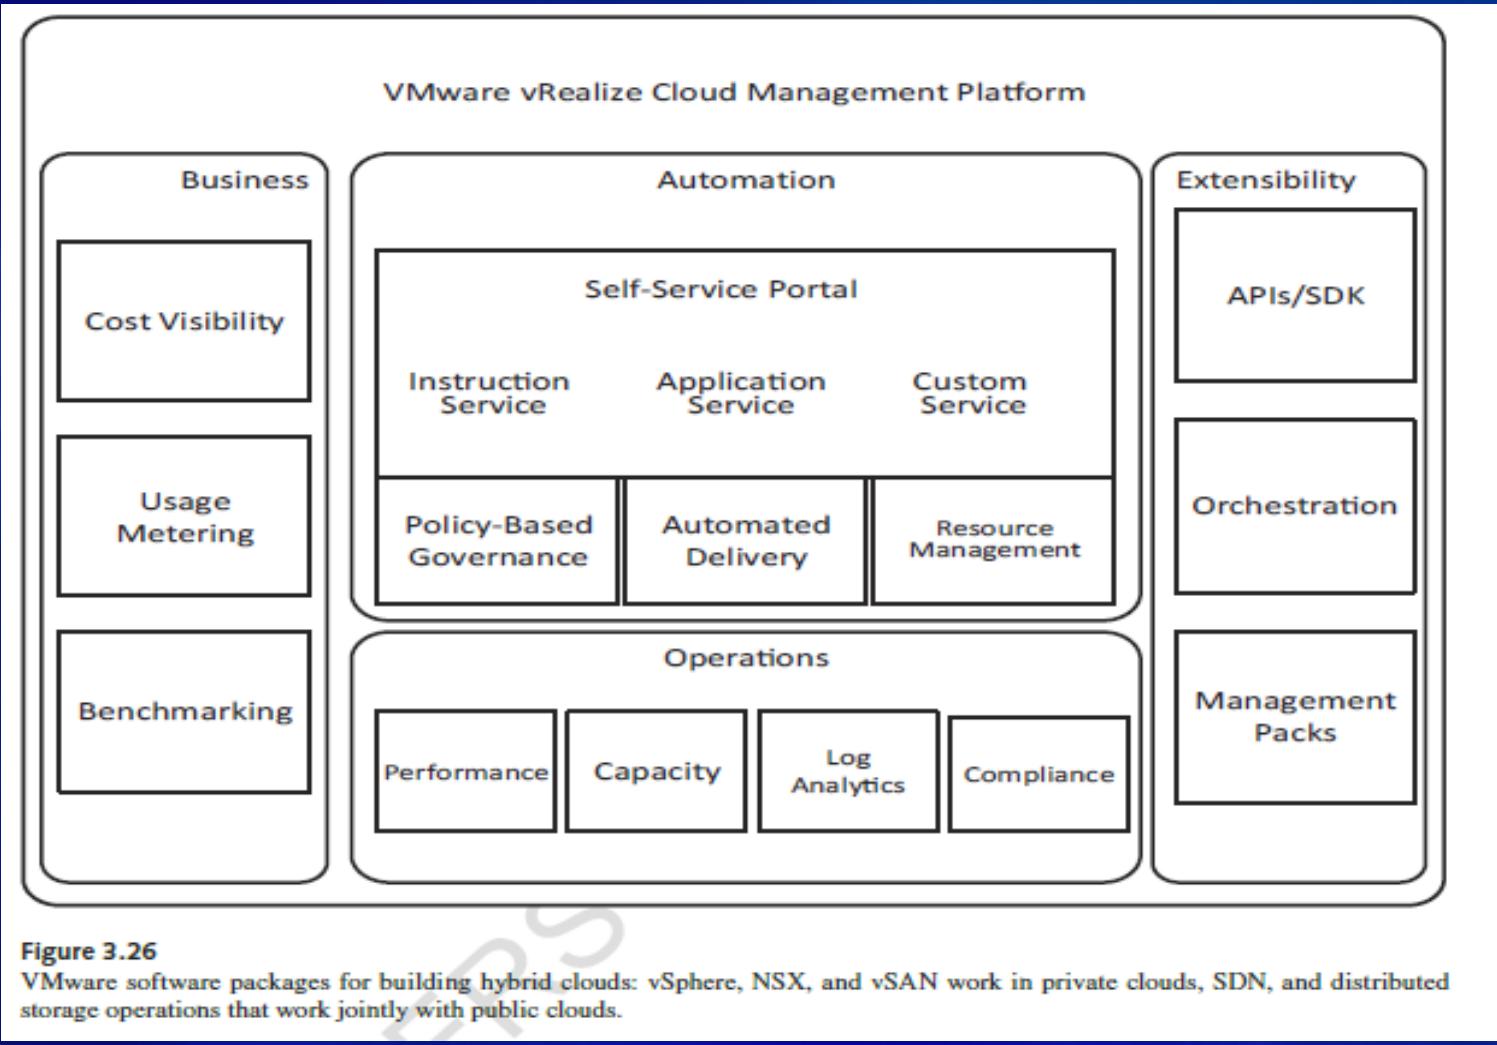
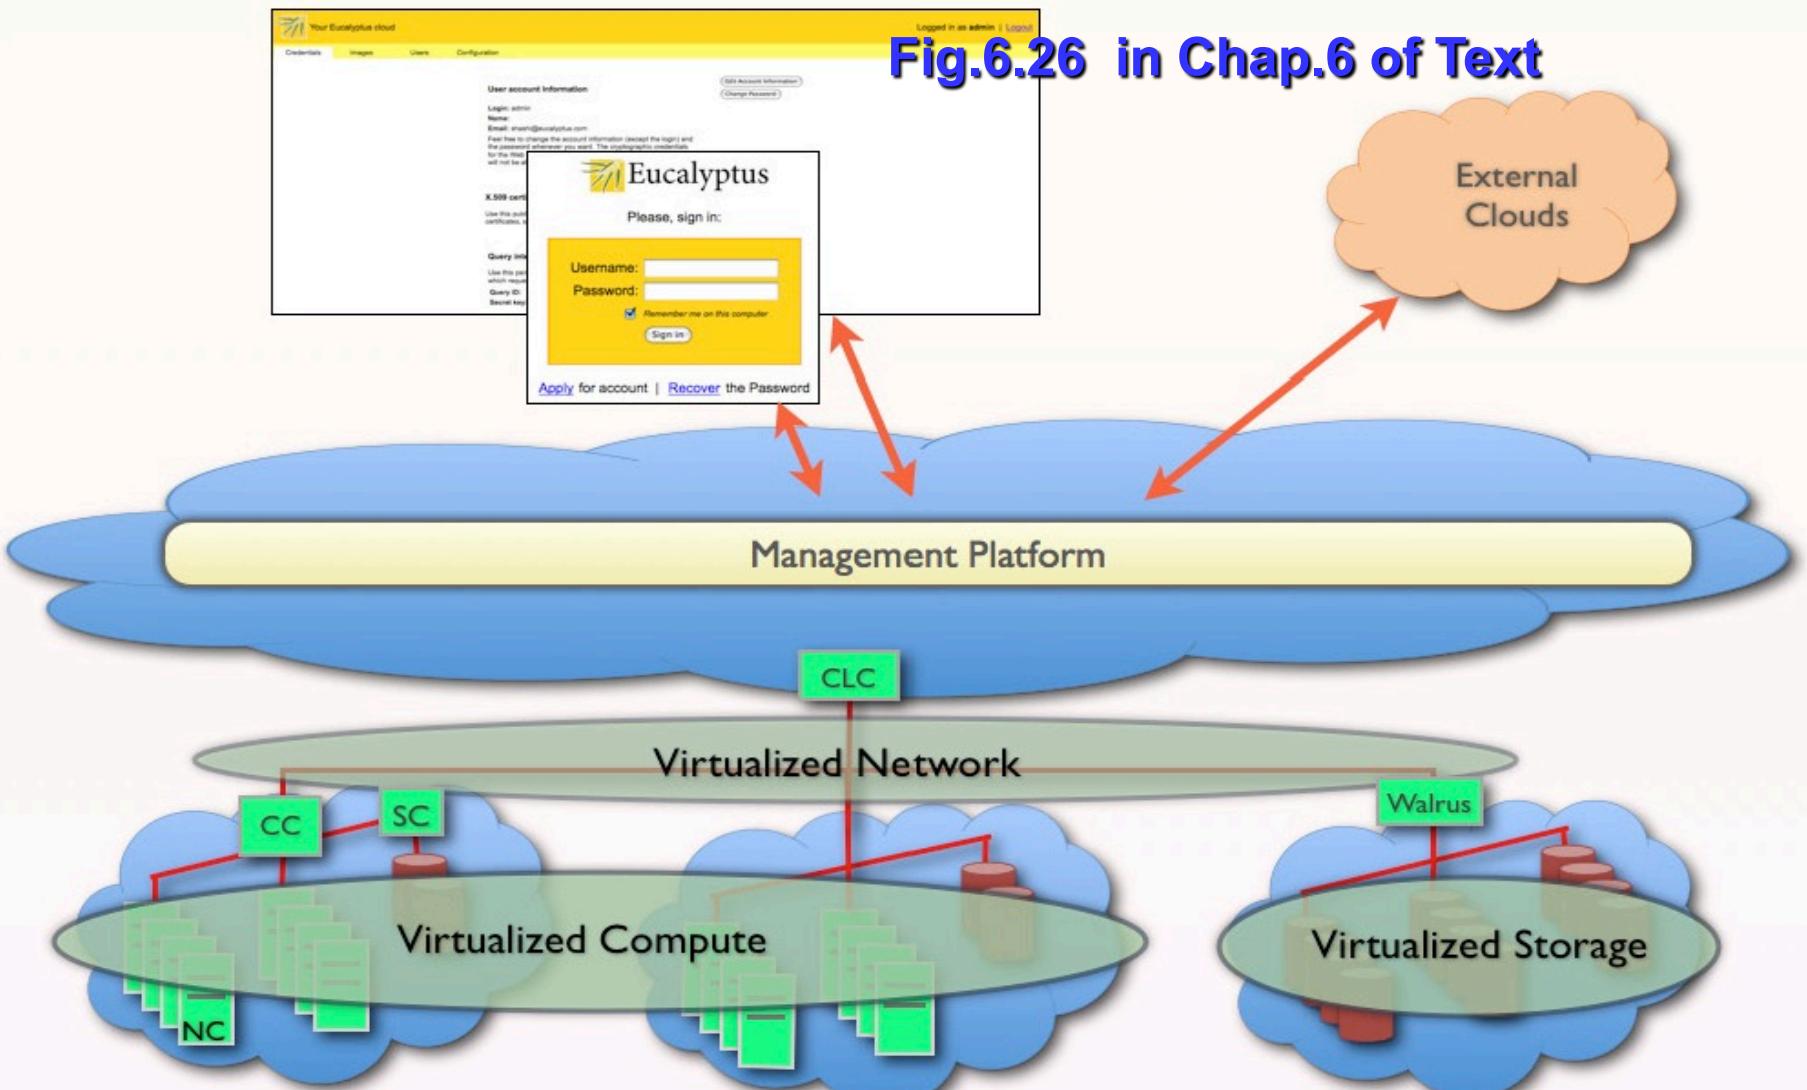


Figure 3.26

VMware software packages for building hybrid clouds: vSphere, NSX, and vSAN work in private clouds, SDN, and distributed storage operations that work jointly with public clouds.

3.17: The vSphere 6 is a cloud OS commercially available from VMware. Review the open literature that reports the porting and application experiences and measured performance by its clients or user groups. Write a short technical report to summarize your research findings.

Eucalyptus architecture



Eucalyptus

An open-source OS for building private clouds (1)

Eucalyptus is an open source software system (Figure 3.27) intended mainly for supporting Infrastructure as a Service (IaaS) clouds. The system primarily supports virtual networking and the management of VMs; virtual storage is not supported. Its purpose is to build private clouds that can interact with end users through Ethernet or the Internet. The system also supports interaction with other private clouds or public clouds over the Internet. The system is short on security and other desired features for general-purpose grid or cloud applications.

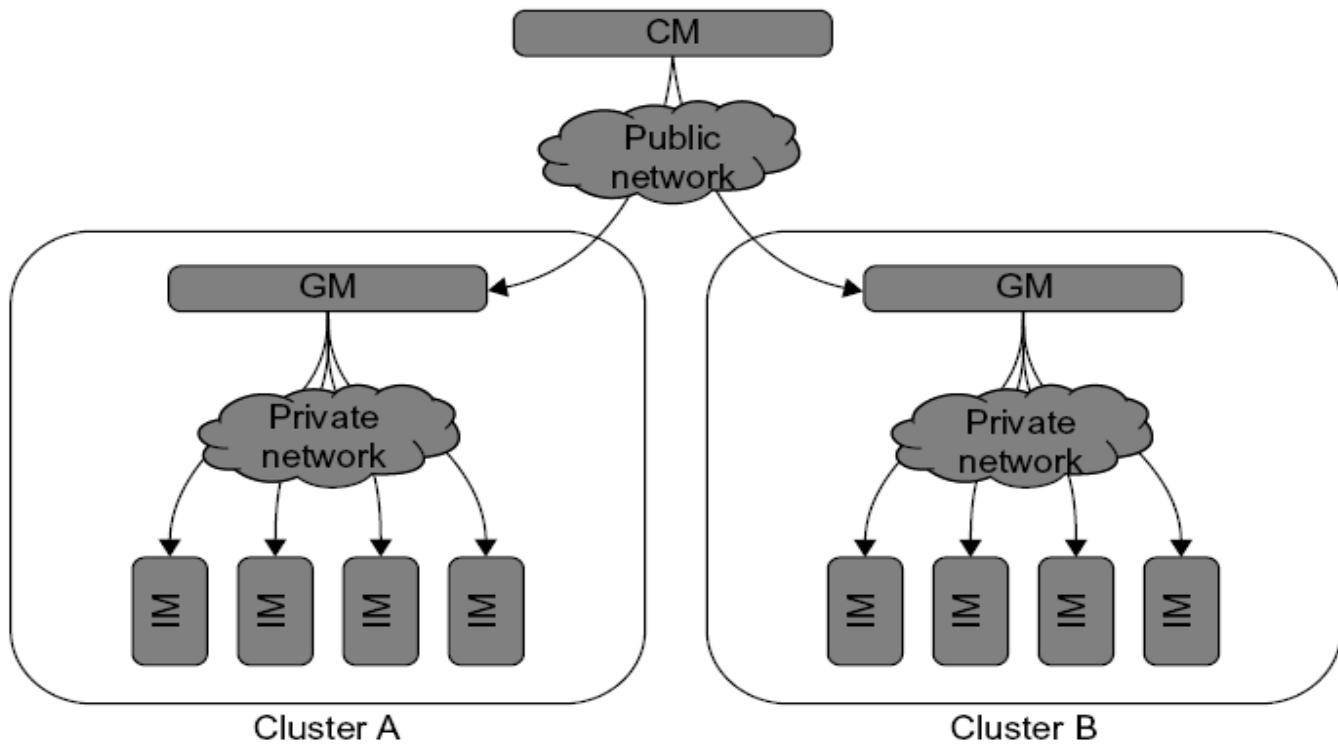


FIGURE 3.27

Eucalyptus for building private clouds by establishing virtual networks over the VMs linking through Ethernet and the Internet.

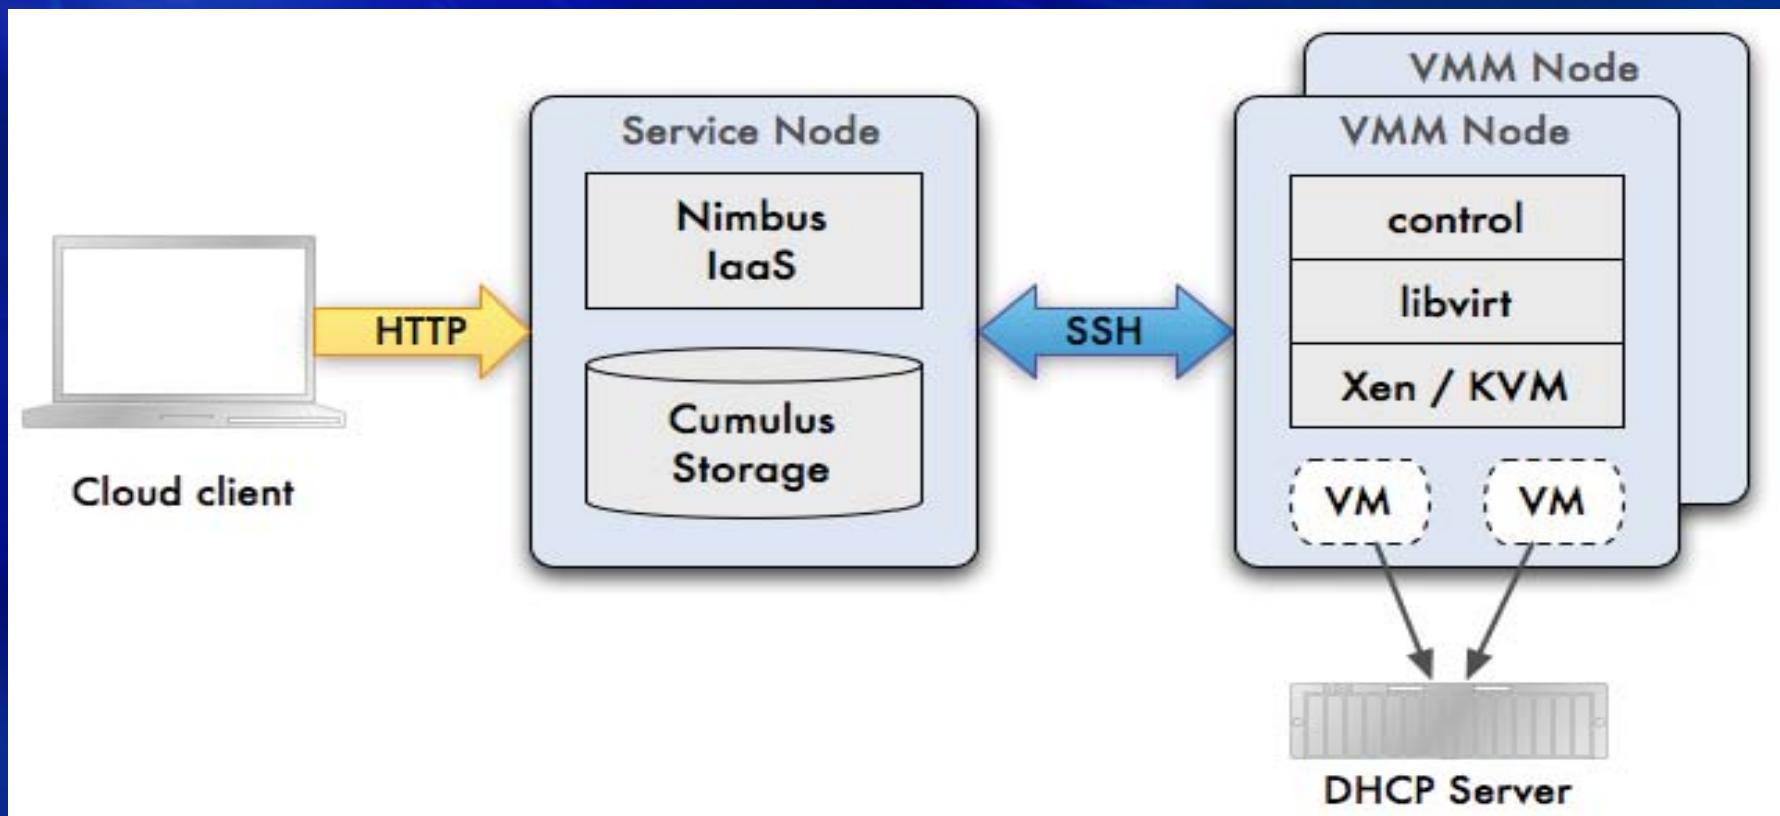
(Courtesy of D. Nurmi, et al. [45])

Eucalyptus

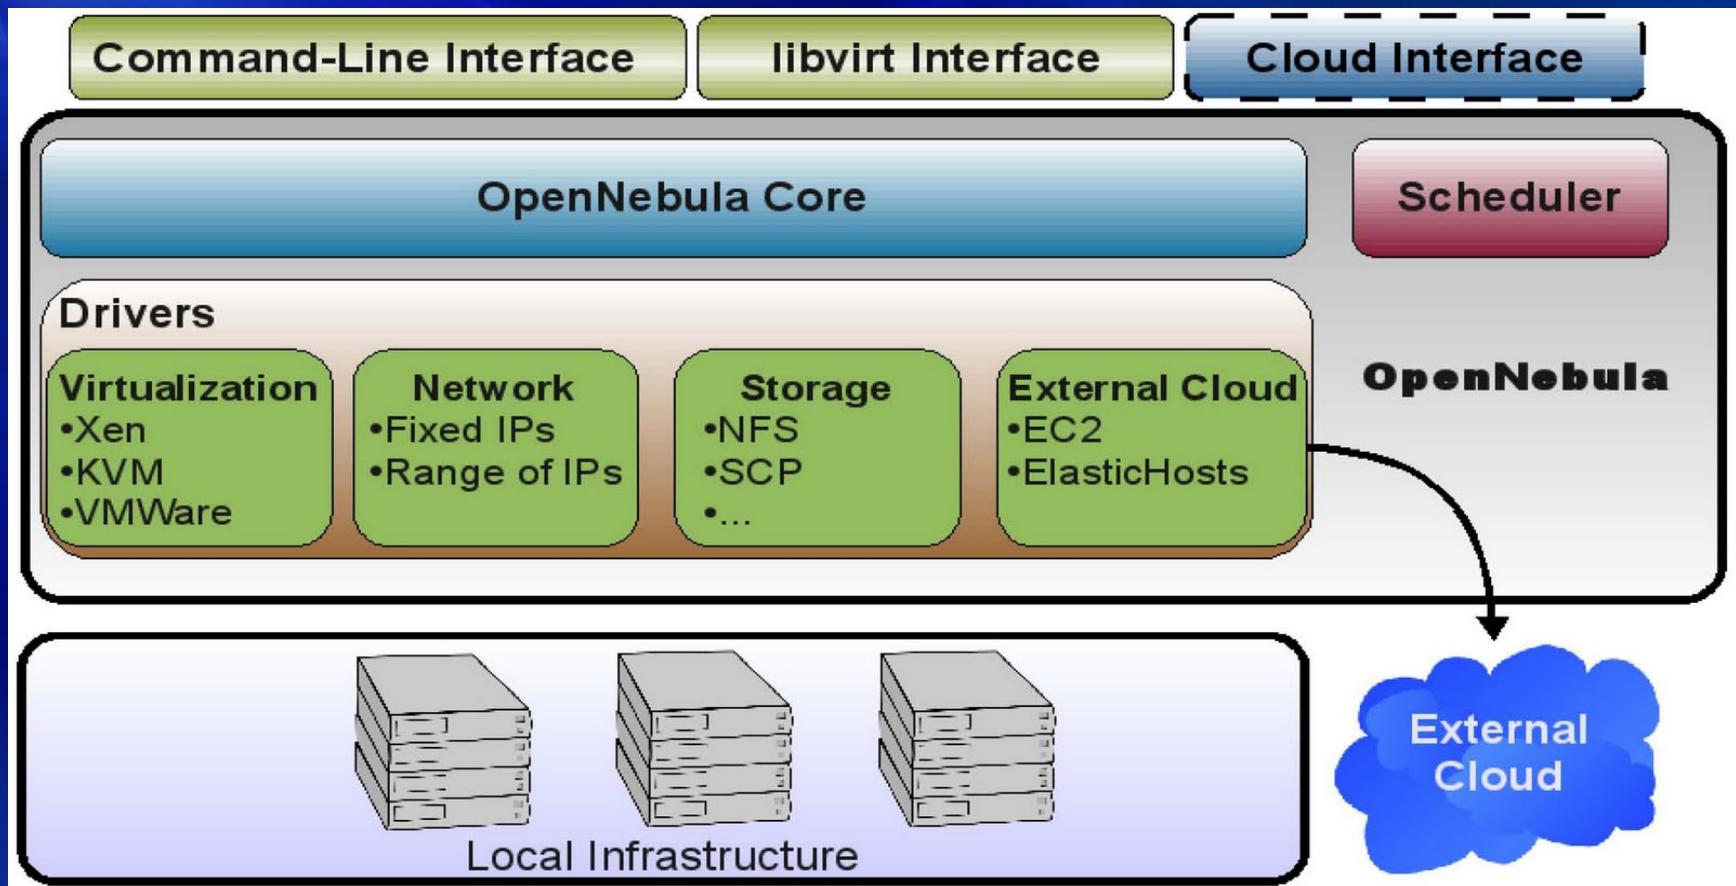
An open-source OS for building private clouds

- **Instance Manager** controls the execution, inspection, and terminating of VM instances on the host where it runs.
- **Group Manager** gathers information about and schedules VM execution on specific instance managers, as well as manages virtual instance network.
- **Cloud Manager** is the entry-point into the cloud for users and administrators. It queries node managers for information about resources, makes scheduling decisions, and implements them by making requests to group managers.

Nimbus cloud infrastructure



OpenNebula architecture



Openstack: An IaaS cloud project launched by Rackspace and NASA in 2010

- Currently, 120 companies have joined Openstack
- Openstack is used to create private cloud and offer cloud computing (Nova), object storage (Swift) and image services (Glance)
- The project offers free open source software under the Apache license. The Openstack cloud software is written in Python: <http://openstack.org/>

Openstack:

An IaaS Cloud Library

- **Nova site:**
<http://openstack.org/projects/compute/>,
[http://launchpad.net/nova/](http://launchpad.net/nova)
- **Swift site:**
<http://openstack.org/projects/storage/>,
[http://launchpad.net/swift/](http://launchpad.net/swift)
- **Glance site:**
<http://openstack.org/projects/image-service/>,
[http://launchpad.net/glance/](http://launchpad.net/glance)

OpenStack Nova system

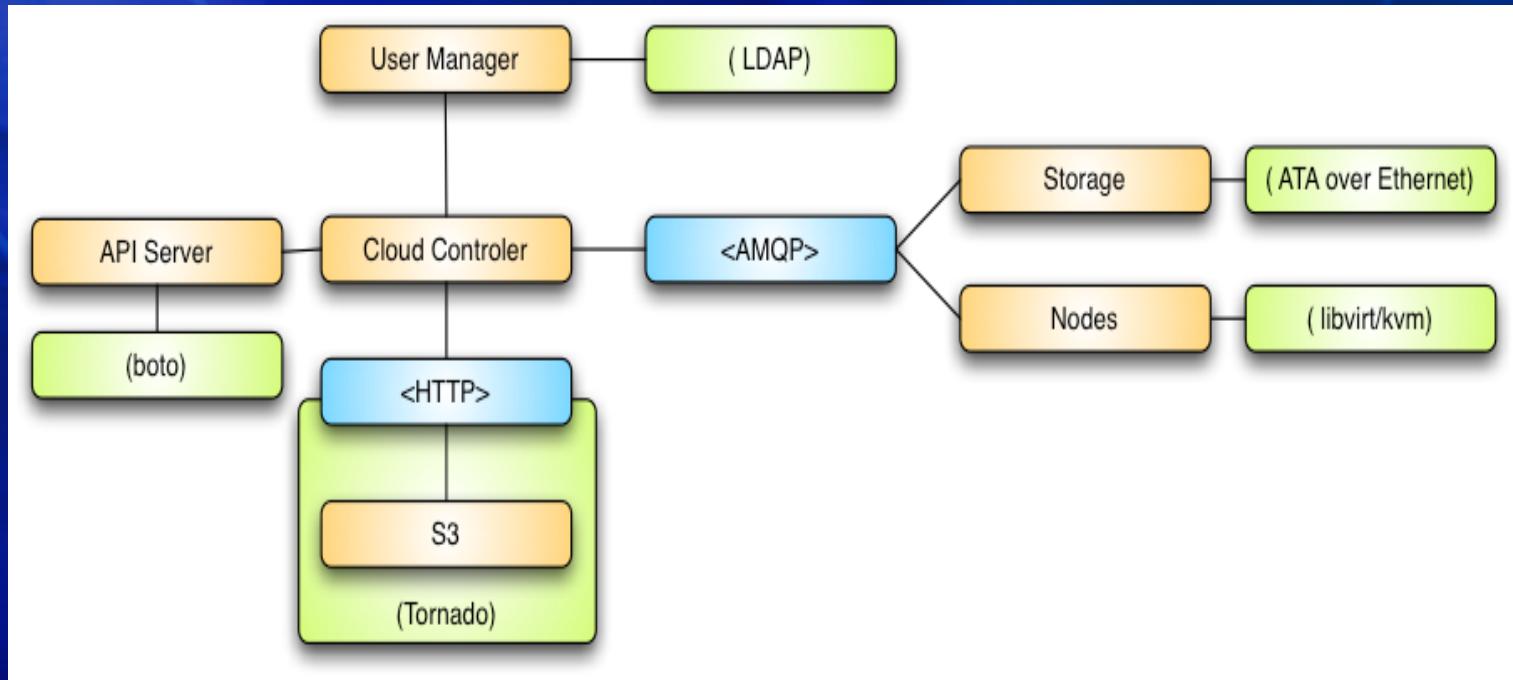


Figure 6.36. OpenStack Nova system architecture. AMQP is an advanced messaging queuing protocol

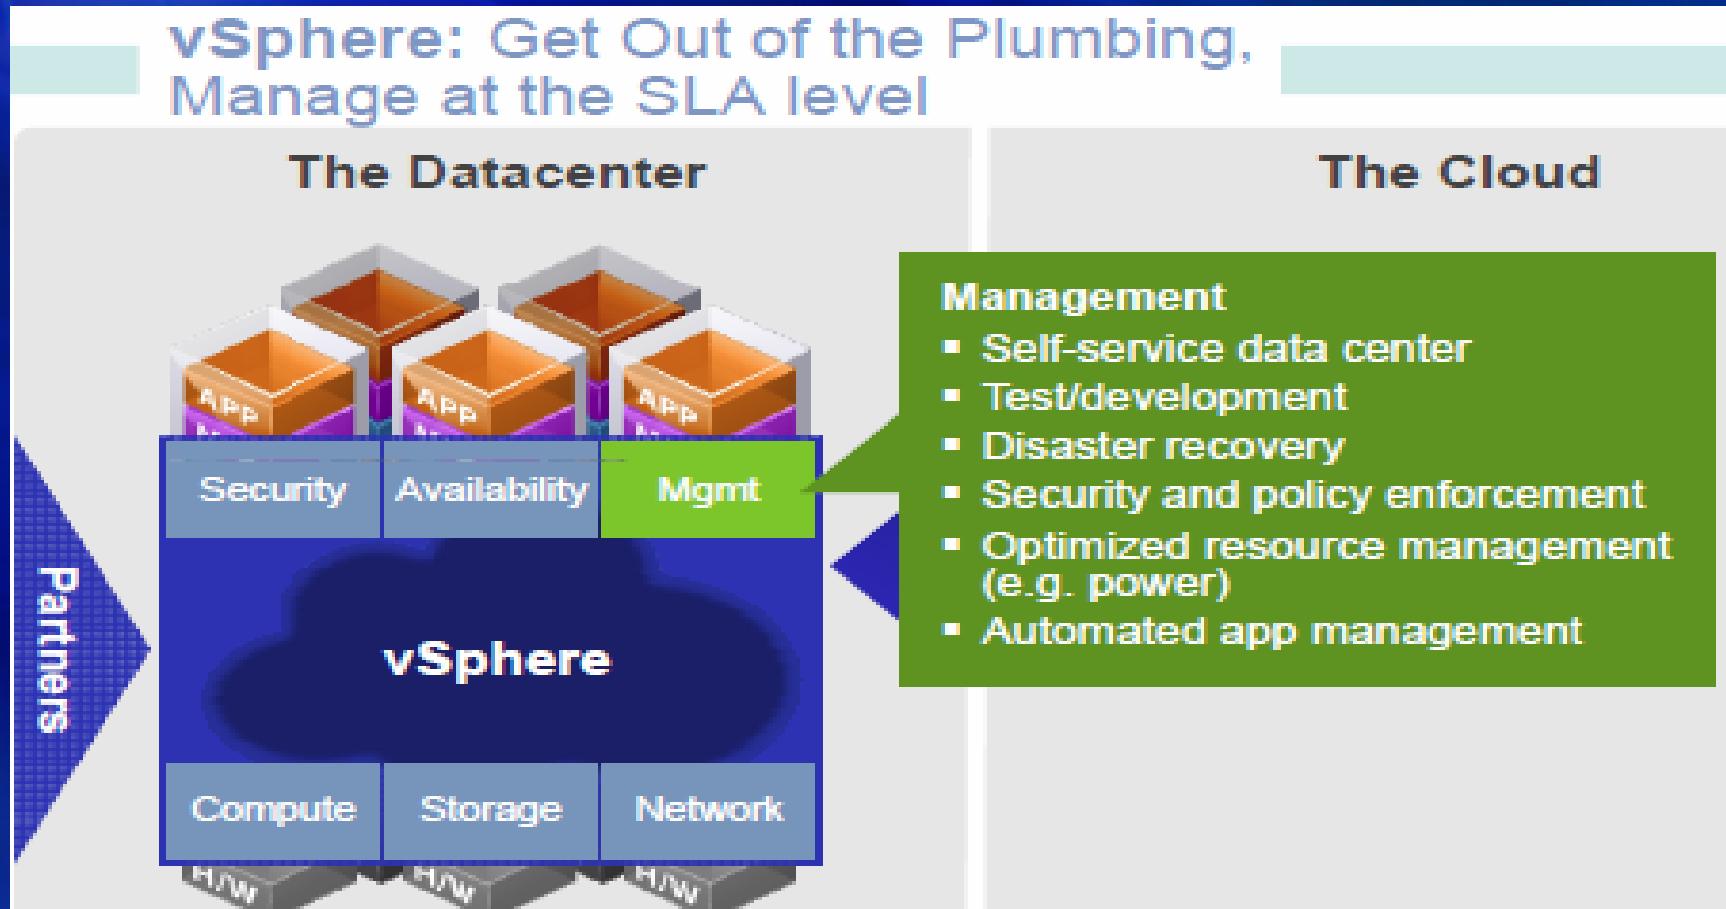
Virtualization Services (HaaS): VMware

- **VMware provides virtualization software with a market share of more than 80%. The company was acquired by EMC in 2004 for \$625 million**
- **VMware Workstation: This software suite allows users to run multiple instances of x86 or x86-64 -compatible OS on a single physical PC**
- **VMware Fusion: This provides similar functionality like the VMware Workstation for users of the Intel Mac platform, along with full compatibility with virtual machines created by other VMware products**

Virtualization services (HaaS): VMware

- **VMware Server** is provided as freeware for non-commercial use, used to create virtual machines with. It is a "hosted" application, which runs within an existing Linux or Windows operating system
- **VMware ESX** is an enterprise-level product that can deliver greater performance than the freeware VMware Server due to lower system overhead. VMware ESX is a "bare-metal" product, running directly on the server hardware, allowing virtual servers to also use hardware more or less directly
- **VMware vSphere** is a "cloud OS" capable of managing large pools of infrastructure. We will cover cloud OS in lectures 15 and 16

vSphere: A commercial OS for converting datacenters into cloud platforms



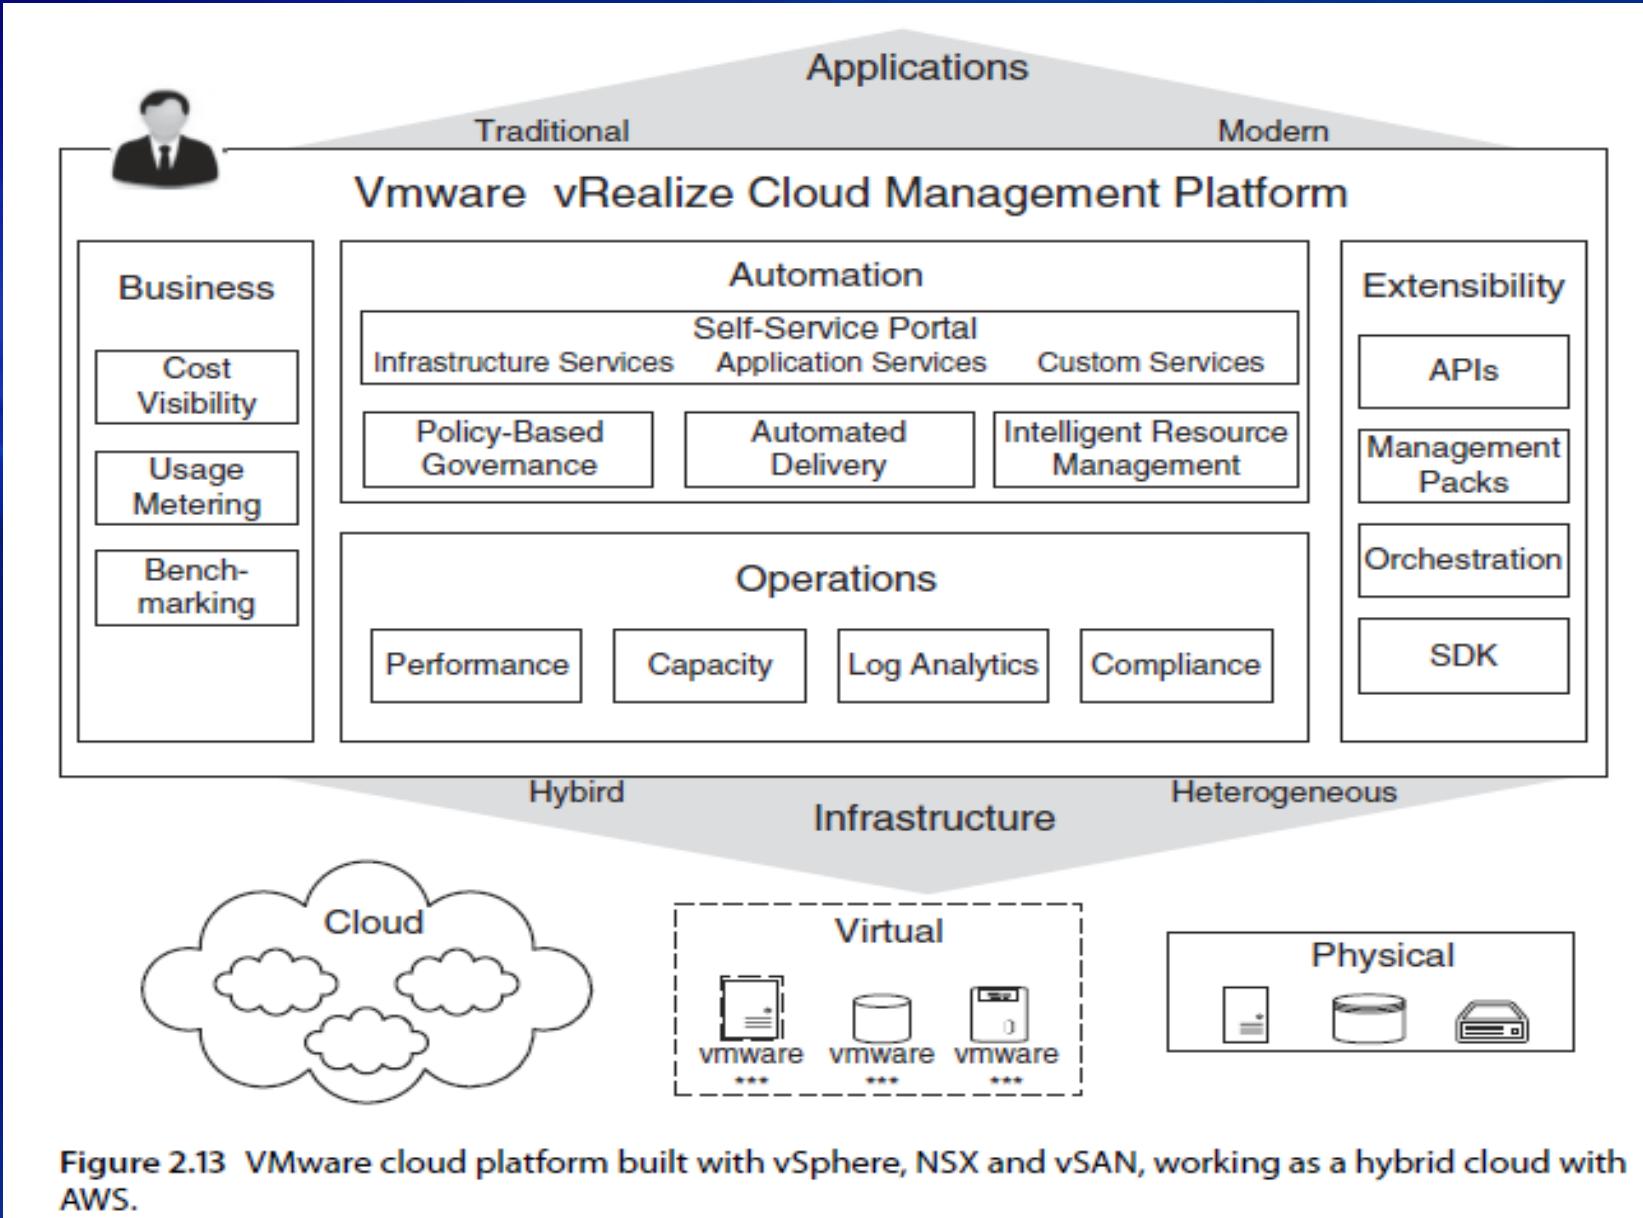


Figure 2.13 VMware cloud platform built with vSphere, NSX and vSAN, working as a hybrid cloud with AWS.

Cloud services: Constraints, cost and productivity

Less Constrained

Constraints in the App Model

More Constrained

Horizontal Software

Compute & Storage



Horizontal Service



Vertical Service
Salesforce



Less Automation

Cost

More Automation

Less

Productivity (assuming a match of app to programming model)

More

Prof. Kai

Advances in HLL for clouds

Table 6.7: Comparison of High Level Data Analysis Languages

	Sawzall	Pig Latin	DryadLINQ
Origin	Google	Yahoo	Microsoft
Data Model	Google Protocol Buffer or basic	Atom, Tuple, Bag, Map	Partition File
Typing	Static	Dynamic	Static
Category	Interpreted	Compiled	Compiled
Programming Style	Imperative	Procedural: sequence of declarative steps	Imperative and Declarative
Similarity to SQL	Least	Moderate	A lot!
Extensibility (User defined functions)	No	Yes	Yes
Control Structures	Yes	No	Yes
Execution Model	Record Operations + fixed aggregations	Sequence of MapReduce operations	Directed Acyclic Graphs
Target Runtime	Google MapReduce	Hadoop (Pig)	Dryad

Reading Assignments

1. K. Hwang, *Cloud Computing for Machine Learning and Cognitive Applications*, Chapter 3, MIT Press, 2017
2. M. Rosenblum and T. Garfinkel, “Virtual Machine Monitors: Current Technology and Future Trends”, *IEEE Computer Magazine*, May 2005, pp. 39-47
3. VM Ware, Inc., *Virtualization Overview*, White paper, <http://www.vmware.com>, 2006
4. J. Smith and R. Nair, *Virtual Machines*, Elsevier, 2005